Red Hat Enterprise Linux 7

Performance Tuning Guide

Optimizing subsystem throughput in Red Hat Enterprise Linux 7

Last Updated: 2017-11-06

Red Hat Enterprise Linux 7 Performance Tuning Guide

Optimizing subsystem throughput in Red Hat Enterprise Linux 7

Milan Navrátil
Red Hat Customer Content Services

Laura Bailey
Red Hat Customer Content Services

Charlie Boyle
Red Hat Customer Content Services

Red Hat Subject Matter Experts

Edited by

Marek Suchánek
Red Hat Customer Content Services
msuchane@redhat.com

Legal Notice

Copyright © 2017 Red Hat, Inc.

This document is licensed by Red Hat under the Creative Commons Attribution-ShareAlike 3.0 Unported License. If you distribute this document, or a modified version of it, you must provide attribution to Red Hat, Inc. and provide a link to the original. If the document is modified, all Red Hat trademarks must be removed.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux ® is the registered trademark of Linus Torvalds in the United States and other countries.

Java ® is a registered trademark of Oracle and/or its affiliates.

XFS ® is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL ® is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js ® is an official trademark of Joyent. Red Hat Software Collections is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack ® Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

Abstract

The Red Hat Enterprise Linux 7 Performance Tuning Guide explains how to optimize Red Hat Enterprise Linux 7 performance. It also documents performance-related upgrades in Red Hat Enterprise Linux 7. The Performance Tuning Guide presents only field-tested and proven procedures. Nonetheless, all prospective configurations should be set up and tested in a testing environment before being applied to a production system. Backing up all data and configuration settings prior to tuning is also recommended. To expand your expertise, you might also be interested in the Red Hat Enterprise Linux Performance Tuning (RH442) training course.

Table of Contents

CHAPTER 1. INTRODUCTION **4**
Who Should Read This Book 4

CHAPTER 2. PERFORMANCE MONITORING TOOLS **5**
2.1. /PROC 5
2.2. GNOME SYSTEM MONITOR 5
2.3. PERFORMANCE CO-PILOT (PCP) 6
2.4. BUILT-IN COMMAND-LINE TOOLS 6
2.5. PERF 7
2.6. TURBOSTAT 8
2.7. IOSTAT 8
2.8. IRQBALANCE 8
2.9. SS 8
2.10. NUMASTAT 9
2.11. NUMAD 9
2.12. SYSTEMTAP 9
2.13. OPROFILE 10
2.14. VALGRIND 10

CHAPTER 3. TUNED **12**
3.1. TUNED OVERVIEW 12
3.2. PERFORMANCE TUNING WITH TUNED AND TUNED-ADM 22

CHAPTER 4. TUNA **25**
4.1. REVIEWING THE SYSTEM WITH TUNA 25
4.2. TUNING CPUS WITH TUNA 26
4.3. TUNING IRQS WITH TUNA 26
4.4. TUNING TASKS WITH TUNA 27
4.5. EXAMPLES OF USING TUNA 28

CHAPTER 5. CPU **30**
5.1. CONSIDERATIONS 30
5.2. MONITORING AND DIAGNOSING PERFORMANCE PROBLEMS 35
5.3. CONFIGURATION SUGGESTIONS 36

CHAPTER 6. MEMORY **42**
6.1. CONSIDERATIONS 42
6.2. MONITORING AND DIAGNOSING PERFORMANCE PROBLEMS 42
6.3. CONFIGURING HUGETLB HUGE PAGES 46
6.4. CONFIGURING TRANSPARENT HUGE PAGES 49
6.5. CONFIGURING SYSTEM MEMORY CAPACITY 49

CHAPTER 7. STORAGE AND FILE SYSTEMS **54**
7.1. CONSIDERATIONS 54
7.2. MONITORING AND DIAGNOSING PERFORMANCE PROBLEMS 59
7.3. SOLID-STATE DISKS 62
7.4. CONFIGURATION TOOLS 63

CHAPTER 8. NETWORKING **75**
8.1. CONSIDERATIONS 75
8.2. MONITORING AND DIAGNOSING PERFORMANCE PROBLEMS 76
8.3. CONFIGURATION TOOLS 77

APPENDIX A. TOOL REFERENCE **84**
A.1. IRQBALANCE 84
A.2. ETHTOOL 85
A.3. SS 85
A.4. TUNED 85
A.5. TUNED-ADM 85
A.6. PERF 87
A.7. PERFORMANCE CO-PILOT (PCP) 88
A.8. VMSTAT 90
A.9. X86_ENERGY_PERF_POLICY 91
A.10. TURBOSTAT 92
A.11. NUMASTAT 93
A.12. NUMACTL 94
A.13. NUMAD 95
A.14. OPROFILE 96
A.15. TASKSET 97
A.16. SYSTEMTAP 98

APPENDIX B. REVISION HISTORY **99**

CHAPTER 1. INTRODUCTION

To learn about features introduced in each Red Hat Enterprise Linux 7 minor release, see the Release Notes for respective minor versions.

The Performance Tuning Guide is a comprehensive guide to optimizing the various sub-systems that make up Red Hat Enterprise Linux 7 for specific purposes. This guide also outlines the performance monitoring and tuning tools available in Red Hat Enterprise Linux 7.

Before you begin tuning, Red Hat has the following important recommendations:

Back up before you configure

The default settings in Red Hat Enterprise Linux 7 are suitable for most services running under moderate loads. Improving the performance of a specific sub-system may affect another system adversely. Back up all data and configuration information before you begin tuning your system.

Test configuration out of production

The procedures documented in the Performance Tuning Guide are tested extensively by Red Hat engineers in both lab and field. Nonetheless, Red Hat recommends testing all planned configurations in a secure testing environment before applying those configurations to production systems.

Who Should Read This Book

The Performance Tuning Guide has been written primarily for two distinct but overlapping audiences:

System administrators

The Performance Tuning Guide documents the effects of each configuration option in detail so that system administrators can optimize Red Hat Enterprise Linux 7 for their specific purpose. Procedures in this guide are suitable for system administrators with Red Hat Certified Engineer (RHCE) certification or an equivalent amount of experience (3–5 years' experience deploying and managing Linux-based systems).

System and business analysts

This guide explains Red Hat Enterprise Linux 7 performance features at a high level. It provides information about how sub-systems perform under specific workloads, allowing analysts to determine whether Red Hat Enterprise Linux 7 is suitable for their use case.

Where possible, the Performance Tuning Guide also refers readers to more detailed feature documentation. This allows readers to develop the in-depth knowledge required to formulate the detailed deployment and optimization strategies necessary for infrastructure and deployment proposals.

CHAPTER 2. PERFORMANCE MONITORING TOOLS

This chapter briefly describes some of the performance monitoring and configuration tools available for Red Hat Enterprise Linux 7. Where possible, this chapter directs readers to further information about how to use the tool, and examples of real life situations that the tool can be used to resolve.

The following knowledge base article provides a more comprehensive list of performance monitoring tools suitable for use with Red Hat Enterprise Linux: https://access.redhat.com/site/solutions/173863.

2.1. /PROC

The **/proc** "file system" is a directory that contains a hierarchy of files that represent the current state of the Linux kernel. It allows users and applications to see the kernel's view of the system.

The **/proc** directory also contains information about system hardware and any currently running processes. Most files in the **/proc** file system are read-only, but some files (primarily those in /proc/sys) can be manipulated by users and applications to communicate configuration changes to the kernel.

For further information about viewing and editing files in the **/proc** directory, refer to the Red Hat Enterprise Linux 7 System Administrator's Reference Guide, available from http://access.redhat.com/site/documentation/Red_Hat_Enterprise_Linux/.

2.2. GNOME SYSTEM MONITOR

The GNOME desktop environment includes a graphical tool, System Monitor, to assist you in monitoring and modifying the behavior of your system. System Monitor displays basic system information and allows you to monitor system processes and resource or file system usage.

System Monitor has four tabs, each of which displays different information about the system.

System

This tab displays basic information about the system's hardware and software.

Processes

This tab displays detailed information about active processes and the relationships between those processes. The processes displayed can be filtered to make certain processes easier to find. This tab also lets you perform some actions on the processes displayed, such as start, stop, kill, and change priority.

Resources

This tab displays the current CPU time usage, memory and swap space usage, and network usage.

File Systems

This tab lists all mounted file systems, and provides some basic information about each, such as the file system type, mount point, and memory usage.

To start System Monitor, press the Super key to enter the Activities Overview, type *System Monitor*, and then press Enter.

For more information about System Monitor, see either the Help menu in the application, or the Red Hat Enterprise Linux 7 *System Administrator's Guide*, available from http://access.redhat.com/site/documentation/Red_Hat_Enterprise_Linux/.

2.3. PERFORMANCE CO-PILOT (PCP)

Red Hat Enterprise Linux 7 introduces support for **Performance Co-Pilot** (**PCP**), a suite of tools, services, and libraries for acquiring, storing, and analyzing system-level performance measurements. Its light-weight distributed architecture makes it particularly well-suited for centralized analysis of complex systems. Performance metrics can be added using the Python, Perl, C++, and C interfaces. Analysis tools can use the client APIs (Python, C++, C) directly, and rich web applications can explore all available performance data using a JSON interface.

The Performance Co-Pilot Collection Daemon (**`pmcd`**) is responsible for collecting performance data on the host system, and various client tools, such as **`pminfo`** or **`pmstat`**, can be used to retrieve, display, archive, and process this data on the same host or over the network. The pcp package provides the command-line tools and underlying functionality. The graphical tool also requires the pcp-gui package.

For a list of system services and tools that are distributed with PCP, see Table A.1, “System Services Distributed with Performance Co-Pilot in Red Hat Enterprise Linux 7” and Table A.2, “Tools Distributed with Performance Co-Pilot in Red Hat Enterprise Linux 7”.

Resources

- The manual page named **PCPIntro** serves as an introduction to Performance Co-Pilot. It provides a list of available tools as well as a description of available configuration options and a list of related manual pages. By default, comprehensive documentation is installed in the **`/usr/share/doc/pcp-doc/`** directory, notably the *Performance Co-Pilot User's and Administrator's Guide* and *Performance Co-Pilot Programmer's Guide*.

- For information on PCP, see the Index of Performance Co-Pilot (PCP) articles, solutions, tutorials and white papers on the Red Hat Customer Portal.

- If you need to determine what PCP tool has the functionality of an older tool you are already familiar with, see the Side-by-side comparison of PCP tools with legacy tools Red Hat Knowledgebase article.

- See the official PCP documentation for an in-depth description of the Performance Co-Pilot and its usage. If you want to start using PCP on Red Hat Enterprise Linux quickly, see the PCP Quick Reference Guide. The official PCP website also contains a list of frequently asked questions.

2.4. BUILT-IN COMMAND-LINE TOOLS

Red Hat Enterprise Linux 7 provides several tools that can be used to monitor your system from the command line, allowing you to monitor your system outside run level 5. This chapter discusses each tool briefly and provides links to further information about where each tool should be used, and how to use them.

2.4.1. top

The top tool, provided by the procps-ng package, gives a dynamic view of the processes in a running system. It can display a variety of information, including a system summary and a list of tasks currently being managed by the Linux kernel. It also has a limited ability to manipulate processes, and to make configuration changes persistent across system restarts.

By default, the processes displayed are ordered according to the percentage of CPU usage, so that you can easily see the processes consuming the most resources. Both the information top displays and its operation are highly configurable to allow you to concentrate on different usage statistics as required.

For detailed information about using top, see the man page:

```
$ man top
```

2.4.2. ps

The ps tool, provided by the procps-ng package, takes a snapshot of a select group of active processes. By default, the group examined is limited to processes that are owned by the current user and associated with the terminal in which ps is run.

ps can provide more detailed information about processes than top, but by default it provides a single snapshot of this data, ordered by process identifier.

For detailed information about using ps, see the man page:

```
$ man ps
```

2.4.3. Virtual Memory Statistics (vmstat)

The Virtual Memory Statistics tool, vmstat, provides instant reports on your system's processes, memory, paging, block input/output, interrupts, and CPU activity. Vmstat lets you set a sampling interval so that you can observe system activity in near-real time.

vmstat is provided by the procps-ng package. For detailed information about using vmstat, see the man page:

```
$ man vmstat
```

2.4.4. System Activity Reporter (sar)

The System Activity Reporter, sar, collects and reports information about system activity that has occurred so far on the current day. The default output displays the current day's CPU usage at 10 minute intervals from the beginning of the day (00:00:00 according to your system clock).

You can also use the **-i** option to set the interval time in seconds, for example, **sar -i 60** tells sar to check CPU usage every minute.

sar is a useful alternative to manually creating periodic reports on system activity with top. It is provided by the sysstat package. For detailed information about using sar, see the man page:

```
$ man sar
```

2.5. PERF

The perf tool uses hardware performance counters and kernel tracepoints to track the impact of other commands and applications on your system. Various perf subcommands display and record statistics for common performance events, and analyze and report on the data recorded.

For detailed information about perf and its subcommands, see Section A.6, “perf”.

Alternatively, more information is available in the Red Hat Enterprise Linux 7 *Developer Guide*, available from http://access.redhat.com/site/documentation/Red_Hat_Enterprise_Linux/.

2.6. TURBOSTAT

Turbostat is provided by the kernel-tools package. It reports on processor topology, frequency, idle power-state statistics, temperature, and power usage on Intel® 64 processors.

Turbostat is useful for identifying servers that are inefficient in terms of power usage or idle time. It also helps to identify the rate of system management interrupts (SMIs) occurring on the system. It can also be used to verify the effects of power management tuning.

Turbostat requires root privileges to run. It also requires processor support for the following:

- invariant time stamp counters
- APERF model-specific registers
- MPERF model-specific registers

For more details about turbostat output and how to read it, see Section A.10, “turbostat”.

For more information about turbostat, see the man page:

```
$ man turbostat
```

2.7. IOSTAT

The **iostat** tool, provided by the sysstat package, monitors and reports on system input/output device loading to help administrators make decisions about how to balance input/output load between physical disks. The **iostat** tool reports on processor or device utilization since **iostat** was last run, or since boot. You can focus the output of these reports on specific devices by using the parameters defined in the iostat(1) manual page. For detailed information on the **`await`** value and what can cause its values to be high, see the following Red Hat Knowledgebase article: What exactly is the meaning of value "await" reported by iostat?

2.8. IRQBALANCE

irqbalance is a command line tool that distributes hardware interrupts across processors to improve system performance. For details about **irqbalance**, see Section A.1, “irqbalance” or the man page:

```
$ man irqbalance
```

2.9. SS

ss is a command-line utility that prints statistical information about sockets, allowing administrators to assess device performance over time. By default, ss lists open non-listening TCP sockets that have established connections, but a number of useful options are provided to help administrators filter out statistics about specific sockets.

Red Hat recommends using ss over netstat in Red Hat Enterprise Linux 7.

One common usage is **`ss -tmpie`** which displays detailed information (including internal information) about TCP sockets, memory usage, and processes using the socket.

ss is provided by the iproute package. For more information, see the man page:

-

```
$ man ss
```

2.10. NUMASTAT

The **numastat** tool displays memory statistics for processes and the operating system on a per-NUMA-node basis.

By default, **numastat** displays per-node NUMA hit an miss system statistics from the kernel memory allocator. Optimal performance is indicated by high **numa_hit** values and low **numa_miss** values. **Numastat** also provides a number of command line options, which can show how system and process memory is distributed across NUMA nodes in the system.

It can be useful to cross-reference per-node **numastat** output with per-CPU **top** output to verify that process threads are running on the same node to which memory is allocated.

Numastat is provided by the numactl package. For details about how to use numastat, see Section A.11, “numastat”. For further information about numastat, see the man page:

```
$ man numastat
```

2.11. NUMAD

numad is an automatic NUMA affinity management daemon. It monitors NUMA topology and resource usage within a system in order to dynamically improve NUMA resource allocation and management (and therefore system performance). Depending on system workload, numad can provide up to 50 percent improvements in performance benchmarks. It also provides a pre-placement advice service that can be queried by various job management systems to provide assistance with the initial binding of CPU and memory resources for their processes.

numad monitors available system resources on a per-node basis by periodically accessing information in the **/proc** file system. It tries to maintain a specified resource usage level, and rebalances resource allocation when necessary by moving processes between NUMA nodes. numad attempts to achieve optimal NUMA performance by localizing and isolating significant processes on a subset of the system's NUMA nodes.

numad primarily benefits systems with long-running processes that consume significant amounts of resources, and are contained in a subset of the total system resources. It may also benefit applications that consume multiple NUMA nodes' worth of resources; however, the benefits provided by numad decrease as the consumed percentage of system resources increases.

numad is unlikely to improve performance when processes run for only a few minutes, or do not consume many resources. Systems with continuous, unpredictable memory access patterns, such as large in-memory databases, are also unlikely to benefit from using numad.

For further information about using numad, see Section 5.3.5, “Automatic NUMA Affinity Management with numad” or Section A.13, “numad”, or refer to the man page:

```
$ man numad
```

2.12. SYSTEMTAP

SystemTap is a tracing and probing tool that lets you monitor and analyze operating system activities, especially kernel activities, in fine detail. It provides information similar to the output of tools like top, ps, netstat, and iostat, but includes additional options for filtering and analyzing collected data.

SystemTap provides a deeper, more precise analysis of system activities and application behavior to allow you to pinpoint system and application bottlenecks.

For more detailed information about SystemTap, see the Red Hat Enterprise Linux 7 *SystemTap Beginner's Guide* and the Red Hat Enterprise Linux 7 *SystemTap TapSet Reference*. Both books are available from http://access.redhat.com/site/documentation/Red_Hat_Enterprise_Linux/.

2.13. OPROFILE

OProfile is a system-wide performance monitoring tool. It uses the processor's dedicated performance monitoring hardware to retrieve information about the kernel and system executables to determine the frequency of certain events, such as when memory is referenced, the number of second-level cache requests, and the number of hardware requests received. OProfile can also be used to determine processor usage, and to determine which applications and services are used most often.

However, OProfile does have several limitations:

- Performance monitoring samples may not be precise. Because the processor may execute instructions out of order, samples can be recorded from a nearby instruction instead of the instruction that triggered the interrupt.
- OProfile expects processes to start and stop multiple times. As such, samples from multiple runs are allowed to accumulate. You may need to clear the sample data from previous runs.
- OProfile focuses on identifying problems with processes limited by CPU access. It is therefore not useful for identifying processes that are sleeping while they wait for locks on other events.

For more detailed information about OProfile, see Section A.14, “OProfile”, or the *Red Hat Enterprise Linux 7 System Administrator's Guide*, available from http://access.redhat.com/site/documentation/Red_Hat_Enterprise_Linux/. Alternatively, refer to the documentation on your system, located in **/usr/share/doc/oprofile-*version***.

2.14. VALGRIND

Valgrind provides a number of detection and profiling tools to help improve the performance of your applications. These tools can detect memory and thread-related errors, as well as heap, stack, and array overruns, letting you easily locate and correct errors in your application code. They can also profile the cache, the heap, and branch-prediction to identify factors that may increase application speed and minimize memory usage.

Valgrind analyzes your application by running it on a synthetic CPU and instrumenting existing application code as it is executed. It then prints commentary that clearly identifies each process involved in application execution to a user-specified file, file descriptor, or network socket. Note that executing instrumented code can take between four and fifty times longer than normal execution.

Valgrind can be used on your application as-is, without recompiling. However, because Valgrind uses debugging information to pinpoint issues in your code, if your application and support libraries were not compiled with debugging information enabled, Red Hat recommends recompiling to include this information.

Valgrind also integrates with the GNU Project Debugger (gdb) to improve debugging efficiency.

Valgrind and its subordinate tools are useful for memory profiling. For detailed information about using Valgrind to profile system memory, see Section 6.2.2, “Profiling Application Memory Usage with Valgrind”.

For detailed information about Valgrind, see the Red Hat Enterprise Linux 7 Developer Guide, available from http://access.redhat.com/site/documentation/Red_Hat_Enterprise_Linux/.

For detailed information about using Valgrind, see the man page:

```
$ man valgrind
```

Accompanying documentation can also be found in **/usr/share/doc/valgrind-*version*** when the valgrind package is installed.

CHAPTER 3. TUNED

3.1. TUNED OVERVIEW

Tuned is a daemon that uses **udev** to monitor connected devices and statically and dynamically tunes system settings according to a selected profile. **Tuned** is distributed with a number of predefined profiles for common use cases like high throughput, low latency, or powersave. It is possible to modify the rules defined for each profile and customize how to tune a particular device. To revert all changes made to the system settings by a certain profile, you can either switch to another profile or deactivate the **tuned** service.

NOTE

Starting with Red Hat Enterprise Linux 7.2, you can run **Tuned** in **`no-daemon mode`**, which does not require any resident memory. In this mode, **tuned** applies the settings and exits. The **`no-daemon mode`** is disabled by default because a lot of **tuned** functionality is missing in this mode, including D-Bus support, hot-plug support, or rollback support for settings. To enable **`no-daemon mode`**, set the following in the **`/etc/tuned/tuned-main.conf`** file: **`daemon = 0`**.

Static tuning mainly consists of the application of predefined **`sysctl`** and **`sysfs`** settings and one-shot activation of several configuration tools like **ethtool**. **Tuned** also monitors the use of system components and tunes system settings dynamically based on that monitoring information.

Dynamic tuning accounts for the way that various system components are used differently throughout the uptime for any given system. For example, the hard drive is used heavily during startup and login, but is barely used later when the user might mainly work with applications such as web browsers or email clients. Similarly, the CPU and network devices are used differently at different times. **Tuned** monitors the activity of these components and reacts to the changes in their use.

As a practical example, consider a typical office workstation. Most of the time, the Ethernet network interface is very inactive. Only a few emails go in and out every once in a while or some web pages might be loaded. For those kinds of loads, the network interface does not have to run at full speed all the time, as it does by default. **Tuned** has a monitoring and tuning plug-in for network devices that can detect this low activity and then automatically lower the speed of that interface, typically resulting in a lower power usage. If the activity on the interface increases for a longer period of time, for example because a DVD image is being downloaded or an email with a large attachment is opened, **tuned** detects this and sets the interface speed to maximum to offer the best performance while the activity level is so high. This principle is used for other plug-ins for CPU and hard disks as well.

Dynamic tuning is globally disabled in Red Hat Enterprise Linux and can be enabled by editing the **`/etc/tuned/tuned-main.conf`** file and changing the **`dynamic_tuning`** flag to **`1`**.

3.1.1. Plug-ins

Tuned uses two types of plugins: *monitoring plugins* and *tuning plugins*. Monitoring plugins are used to get information from a running system. Currently, the following monitoring plugins are implemented:

`disk`

Gets disk load (number of IO operations) per device and measurement interval.

`net`

Gets network load (number of transferred packets) per network card and measurement interval.

`load`

Gets CPU load per CPU and measurement interval.

The output of the monitoring plugins can be used by tuning plugins for dynamic tuning. Currently implemented dynamic tuning algorithms try to balance the performance and powersave and are therefore disabled in the performance profiles (dynamic tuning for individual plugins can be enabled or disabled in the **tuned** profiles). Monitoring plugins are automatically instantiated whenever their metrics are needed by any of the enabled tuning plugins. If two tuning plugins require the same data, only one instance of the monitoring plugin is created and the data is shared.

Each tuning plugin tunes an individual subsystem and takes several parameters that are populated from the **tuned** profiles. Each subsystem can have multiple devices (for example, multiple CPUs or network cards) that are handled by individual instances of the tuning plugins. Specific settings for individual devices are also supported. The supplied profiles use wildcards to match all devices of individual subsystems (for details on how to change this, refer to Section 3.1.3, “Custom Profiles”), which allows the plugins to tune these subsystems according to the required goal (selected profile) and the only thing that the user needs to do is to select the correct **tuned** profile.

Currently, the following tuning plugins are implemented (only some of these plugins implement dynamic tuning, parameters supported by plugins are also listed):

`cpu`

Sets the CPU governor to the value specified by the ***`governor`*** parameter and dynamically changes the PM QoS CPU DMA latency according to the CPU load. If the CPU load is lower than the value specified by the ***`load_threshold`*** parameter, the latency is set to the value specified by the ***`latency_high`*** parameter, otherwise it is set to value specified by ***`latency_low`***. Also the latency can be forced to a specific value without being dynamically changed further. This can be accomplished by setting the ***`force_latency`*** parameter to the required latency value.

`eeepc_she`

Dynamically sets the FSB speed according to the CPU load; this feature can be found on some netbooks and is also known as the Asus Super Hybrid Engine. If the CPU load is lower or equal to the value specified by the ***`load_threshold_powersave`*** parameter, the plugin sets the FSB speed to the value specified by the ***`she_powersave`*** parameter (for details about the FSB frequencies and corresponding values, see the kernel documentation, the provided defaults should work for most users). If the CPU load is higher or equal to the value specified by the ***`load_threshold_normal`*** parameter, it sets the FSB speed to the value specified by the ***`she_normal`*** parameter. Static tuning is not supported and the plugin is transparently disabled if the hardware support for this feature is not detected.

`net`

Configures wake-on-lan to the values specified by the ***`wake_on_lan`*** parameter (it uses same syntax as the **ethtool** utility). It also dynamically changes the interface speed according to the interface utilization.

`sysctl`

Sets various **`sysctl`** settings specified by the plugin parameters. The syntax is **`name`**=**`value`**, where **`name`** is the same as the name provided by the **sysctl** tool. Use this plugin if you need to change settings that are not covered by other plugins (but prefer specific plugins if the settings are covered

by them).

usb

Sets autosuspend timeout of USB devices to the value specified by the ***autosuspend*** parameter. The value 0 means that autosuspend is disabled.

vm

Enables or disables transparent huge pages depending on the Boolean value of the ***transparent_hugepages*** parameter.

audio

Sets the autosuspend timeout for audio codecs to the value specified by the ***timeout*** parameter. Currently **snd_hda_intel** and **snd_ac97_codec** are supported. The value **0** means that the autosuspend is disabled. You can also enforce the controller reset by setting the Boolean parameter ***reset_controller*** to **true**.

disk

Sets the elevator to the value specified by the ***elevator*** parameter. It also sets ALPM to the value specified by the ***alpm*** parameter, ASPM to the value specified by the ***aspm*** parameter, scheduler quantum to the value specified by the ***scheduler_quantum*** parameter, disk spindown timeout to the value specified by the ***spindown*** parameter, disk readahead to the value specified by the ***readahead*** parameter, and can multiply the current disk readahead value by the constant specified by the ***readahead_multiply*** parameter. In addition, this plugin dynamically changes the advanced power management and spindown timeout setting for the drive according to the current drive utilization. The dynamic tuning can be controlled by the Boolean parameter ***dynamic*** and is enabled by default.

mounts

Enables or disables barriers for mounts according to the Boolean value of the ***disable_barriers*** parameter.

script

This plugin can be used for the execution of an external script that is run when the profile is loaded or unloaded. The script is called by one argument which can be **start** or **stop** (it depends on whether the script is called during the profile load or unload). The script file name can be specified by the ***script*** parameter. Note that you need to correctly implement the stop action in your script and revert all setting you changed during the start action, otherwise the roll-back will not work. For your convenience, the **functions** Bash helper script is installed by default and allows you to import and use various functions defined in it. Note that this functionality is provided mainly for backwards compatibility and it is recommended that you use it as the last resort and prefer other plugins if they cover the required settings.

sysfs

Sets various **sysfs** settings specified by the plugin parameters. The syntax is **name**=**value**, where **name** is the **sysfs** path to use. Use this plugin in case you need to change some settings that are not covered by other plugins (please prefer specific plugins if they cover the required settings).

video

Sets various powersave levels on video cards (currently only the Radeon cards are supported). The

powersave level can be specified by using the ***radeon_powersave*** parameter. Supported values are: **default**, **auto**, **low**, **mid**, **high**, and **dynpm**. For details, refer to http://www.x.org/wiki/RadeonFeature#KMS_Power_Management_Options. Note that this plugin is experimental and the parameter may change in the future releases.

bootloader

Adds paramaters to the kernel boot command line. This plugin supports the grub-legacy and grub2 packages and also Grub with Extensible Firmware Interface (EFI). Customized non-standard location of the grub2 configuration file can be specified by the ***grub2_cfg_file*** option. The parameters are added to the current grub configuration and its templates. The machine needs to be rebooted for the kernel parameters to take effect.

The parameters can be specified by the following syntax:

```
cmdline=arg1 arg2 ... argn.
```

3.1.2. Installation and Usage

To install the tuned package, run, as root, the following command:

```
yum install tuned
```

Installation of the tuned package also presets the profile which should be the best for you system. Currently the default profile is selected according the following customizable rules:

throughput-performance

This is pre-selected on Red Hat Enterprise Linux 7 operating systems which act as compute nodes. The goal on such systems is the best throughput performance.

virtual-guest

This is pre-selected on virtual machines. The goal is best performance. If you are not interested in best performance, you would probably like to change it to the **balanced** or **powersave** profile (see bellow).

balanced

This is pre-selected in all other cases. The goal is balanced performance and power consumption.

To start **tuned**, run, as root, the following command:

```
systemctl start tuned
```

To enable **tuned** to start every time the machine boots, type the following command:

```
systemctl enable tuned
```

For other **tuned** control such as selection of profiles and other, use:

```
tuned-adm
```

This command requires the **tuned** service to be running.

To view the available installed profiles, run:

```
tuned-adm list
```

To view the currently activated profile, run:

```
tuned-adm active
```

To select or activate a profile, run:

```
tuned-adm profile profile
```

For example:

```
tuned-adm profile powersave
```

As an experimental feature it is possible to select more profiles at once. The **tuned** application will try to merge them during the load. If there are conflicts the settings from the last specified profile will take precedence. This is done automatically and there is no checking whether the resulting combination of parameters makes sense. If used without thinking, the feature may tune some parameters the opposite way which may be counterproductive. An example of such situation would be setting the disk for the **high** throughput by using the **throughput-performance** profile and concurrently setting the disk spindown to the **low** value by the **spindown-disk** profile. The following example optimizes the system for run in a virtual machine for the best performance and concurrently tune it for the low power consumption while the low power consumption is the priority:

```
tuned-adm profile virtual-guest powersave
```

To let **tuned** recommend you the best suitable profile for your system without changing any existing profiles and using the same logic as used during the installation, run the following command:

```
tuned-adm recommend
```

Tuned itself has additional options that you can use when you run it manually. However, this is not recommended and is mostly intended for debugging purposes. The available options can be viewing using the following command:

```
tuned --help
```

3.1.3. Custom Profiles

Distribution-specific profiles are stored in the **/usr/lib/tuned/** directory. Each profile has its own directory. The profile consists of the main configuration file called **tuned.conf**, and optionally other files, for example helper scripts.

If you need to customize a profile, copy the profile directory into the **/etc/tuned/** directory, which is used for custom profiles. If there are two profiles of the same name, the profile included in **/etc/tuned/** is used.

You can also create your own profile in the **/etc/tuned/** directory to use a profile included in **/usr/lib/tuned/** with only certain parameters adjusted or overridden.

The **tuned.conf** file contains several sections. There is one **[main]** section. The other sections are configurations for plugins instances. All sections are optional including the **[main]** section. Lines starting with the hash sign (#) are comments.

The **[main]** section has the following option:

include=*profile*

The specified profile will be included, e.g. **include=powersave** will include the **powersave** profile.

Sections describing plugins instances are formatted in the following way:

```
[NAME]
type=TYPE
devices=DEVICES
```

NAME is the name of the plugin instance as it is used in the logs. It can be an arbitrary string. *TYPE* is the type of the tuning plugin. For a list and descriptions of the tuning plugins refer to Section 3.1.1, “Plug-ins”. *DEVICES* is the list of devices this plugin instance will handle. The **devices** line can contain a list, a wildcard (*), and negation (!). You can also combine rules. If there is no **devices** line all devices present or later attached on the system of the **TYPE** will be handled by the plugin instance. This is same as using **devices=***. If no instance of the plugin is specified, the plugin will not be enabled. If the plugin supports more options, they can be also specified in the plugin section. If the option is not specified, the default value will be used (if not previously specified in the included plugin). For the list of plugin options refer to Section 3.1.1, “Plug-ins”).

Example 3.1. Describing Plug-ins Instances

The following example will match everything starting with **sd**, such as **sda** or **sdb** and doesn't disable barriers on them:

```
[data_disk]
type=disk
devices=sd*
disable_barriers=false
```

The following example will match everything except **sda1** and **sda2**:

```
[data_disk]
type=disk
devices=!sda1, !sda2
disable_barriers=false
```

In cases where you do not need custom names for the plugin instance and there is only one definition of the instance in your configuration file, Tuned supports the following short syntax:

```
[TYPE]
devices=DEVICES
```

In this case, it is possible to omit the **type** line. The instance will then be referred to with a name, same as the type. The previous example could be then rewritten into:

```
[disk]
devices=sdb*
disable_barriers=false
```

If the same section is specified more than once using the **include** option, then the settings are merged. If they cannot be merged due to a conflict, the last conflicting definition overrides the previous settings in conflict. Sometimes, you do not know what was previously defined. In such cases, you can use the **replace** boolean option and set it to **true**. This will cause all the previous definitions with the same name to be overwritten and the merge will not happen.

You can also disable the plugin by specifying the **enabled=false** option. This has the same effect as if the instance was never defined. Disabling the plugin can be useful if you are redefining the previous definition from the **include** option and do not want the plugin to be active in your custom profile.

The following is an example of a custom profile that is based on the **balanced** profile and extends it the way that ALPM for all devices is set to the maximal powersaving.

```
[main]
include=balanced

[disk]
alpm=min_power
```

The following is an example of a custom profile that adds **isolcpus=2** to the kernel boot command line:

```
[bootloader]
cmdline=isolcpus=2
```

The machine needs to be rebooted after the profile is applied for the changes to take effect.

3.1.4. Tuned-adm

A detailed analysis of a system can be very time-consuming. Red Hat Enterprise Linux 7 includes a number of predefined profiles for typical use cases that you can easily activate with the **tuned-adm** utility. You can also create, modify, and delete profiles.

To list all available profiles and identify the current active profile, run:

```
tuned-adm list
```

To only display the currently active profile, run:

```
tuned-adm active
```

To switch to one of the available profiles, run:

```
tuned-adm profile profile_name
```

for example:

```
tuned-adm profile latency-performance
```

To disable all tuning:

```
tuned-adm off
```

The following is a list of profiles which are installed with the base package:

balanced

The default power-saving profile. It is intended to be a compromise between performance and power consumption. It tries to use auto-scaling and auto-tunning whenever possible. It has good results for most loads. The only drawback is the increased latency. In the current **tuned** release it enables the CPU, disk, audio and video plugins and activates the **ondemand** governor. The ***radeon_powersave*** is set to **auto**.

powersave

A profile for maximum power saving performance. It can throttle the performance in order to minimize the actual power consumption. In the current **tuned** release it enables USB autosuspend, WiFi power saving and ALPM power savings for SATA host adapters. It also schedules multi-core power savings for systems with a low wakeup rate and activates the **ondemand** governor. It enables AC97 audio power saving or, depending on your system, HDA-Intel power savings with a 10 seconds timeout. In case your system contains supported Radeon graphics card with enabled KMS it configures it to automatic power saving. On Asus Eee PCs a dynamic Super Hybrid Engine is enabled.

NOTE

The **powersave** profile may not always be the most efficient. Consider there is a defined amount of work that needs to be done, for example a video file that needs to be transcoded. Your machine can consume less energy if the transcoding is done on the full power, because the task will be finished quickly, the machine will start to idle and can automatically step-down to very efficient power save modes. On the other hand if you transcode the file with a throttled machine, the machine will consume less power during the transcoding, but the process will take longer and the overall consumed energy can be higher. That is why the **balanced** profile can be generally a better option.

throughput-performance

A server profile optimized for high throughput. It disables power savings mechanisms and enables sysctl settings that improve the throughput performance of the disk, network IO and switched to the **deadline** scheduler. CPU governor is set to **performance**.

latency-performance

A server profile optimized for low latency. It disables power savings mechanisms and enables sysctl settings that improve the latency. CPU governor is set to **performance** and the CPU is locked to the low C states (by PM QoS).

network-latency

A profile for low latency network tuning. It is based on the **latency-performance** profile. It additionally disables transparent hugepages, NUMA balancing and tunes several other network related sysctl parameters.

network-throughput

Profile for throughput network tuning. It is based on the **throughput-performance** profile. It additionally increases kernel network buffers.

virtual-guest

A profile designed for virtual guests based on the enterprise-storage profile that, among other tasks, decreases virtual memory swappiness and increases disk readahead values. It does not disable disk barriers.

virtual-host

A profile designed for virtual hosts based on the **enterprise-storage** profile that, among other tasks, decreases virtual memory swappiness, increases disk readahead values and enables more aggresive \ of dirty pages.

oracle

A profile optimized for Oracle databases loads based on **throughput-performance** profile. It additionaly disables transparent huge pages and modifies some other performance related kernel parameters. This profile is provided by tuned-profiles-oracle package. It is available in Red Hat Enterprise Linux 6.8 and later.

desktop

A profile optimized for desktops, based on the **balanced** profile. It additionally enables scheduler autogroups for better response of interactive applications.

NOTE

There may be more product specific or 3rd party Tuned profiles available. Such profiles are usually provided by separate RPM packages.

Additional predefined profiles can be installed with the tuned-profiles-compat package available in the **Optional** channel. These profiles are intended for backward compatibility and are no longer developed. The generalized profiles from the base package will mostly perform the same or better. If you do not have specific reason for using them, please prefer the above mentioned profiles from the base package. The compat profiles are following:

default

This has the lowest impact on power saving of the available profiles and only enables CPU and disk plugins of **tuned**.

desktop-powersave

A power-saving profile directed at desktop systems. Enables ALPM power saving for SATA host adapters as well as the CPU, Ethernet, and disk plugins of **tuned**.

laptop-ac-powersave

A medium-impact power-saving profile directed at laptops running on AC. Enables ALPM powersaving for SATA host adapters, Wi-Fi power saving, as well as the CPU, Ethernet, and disk plugins of **tuned**.

laptop-battery-powersave

A high-impact power-saving profile directed at laptops running on battery. In the current **tuned** implementation it is an alias for the **powersave** profile.

spindown-disk

A power-saving profile for machines with classic HDDs to maximize spindown time. It disables the **tuned** power savings mechanism, disables USB autosuspend, disables Bluetooth, enables Wi-Fi power saving, disables logs syncing, increases disk write-back time, and lowers disk swappiness. All partitions are remounted with the **noatime** option.

enterprise-storage

A server profile directed at enterprise-class storage, maximizing I/O throughput. It activates the same settings as the **throughput-performance** profile, multiplies readahead settings, and disables barriers on non-root and non-boot partitions.

NOTE

Use the **atomic-host** profile on physical machines, and the **atomic-guest** profile on virtual machines.

To enable the **tuned** profiles for Red Hat Enterprise Linux Atomic Host, install the tuned-profiles-atomic package. Run, as root, the following command:

```
yum install tuned-profiles-atomic
```

The two **tuned** profiles for Red Hat Enterprise Linux Atomic Host are:

atomic-host

A profile optimized for Red Hat Enterprise Linux Atomic Host, when used as a host system on a bare-metal server, using the throughput-performance profile. It additionally increases SELinux AVC cache, PID limit, and tunes netfilter connections tracking.

atomic-guest

A profile optimized for Red Hat Enterprise Linux Atomic Host, when used as a guest system based on the virtual-guest profile. It additionally increases SELinux AVC cache, PID limit, and tunes netfilter connections tracking.

NOTE

There may be more product-specific or 3rd-party **tuned** profiles available. These profiles are usually provided by separate RPM packages. Three **tuned** profiles are available that enable to edit the kernel command line: **realtime**, **realtime-virtual-host** and **realtime-virtual-guest**.

To enable the **realtime** profile, install the tuned-profiles-realtime package. Run, as root, the following command:

```
yum install tuned-profiles-realtime
```

To enable the **realtime-virtual-host** and **realtime-virtual-guest** profiles, install the tuned-profiles-nfv package. Run, as root, the following command:

```
yum install tuned-profiles-nfv
```

3.1.5. powertop2tuned

The **powertop2tuned** utility is a tool that allows you to create custom **tuned** profiles from the **PowerTOP** suggestions.

To install the **powertop2tuned** application, run the following command as root:

```
yum install tuned-utils
```

To create a custom profile, run the following command as root:

```
powertop2tuned new_profile_name
```

By default it creates the profile in the **/etc/tuned** directory and it bases it on the currently selected **tuned** profile. For safety reasons all **PowerTOP** tunings are initially disabled in the new profile. To enable them uncomment the tunings of your interest in the **/etc/tuned/*profile*/tuned.conf**. You can use the **--enable** or **-e** option that will generate the new profile with most of the tunings suggested by **PowerTOP** enabled. Some dangerous tunings like the USB autosuspend will still be disabled. If you really need them you will have to uncomment them manually. By default, the new profile is not activated. To activate it run the following command:

```
tuned-adm profile new_profile_name
```

For a complete list of the options **powertop2tuned** supports, type in the following command:

```
powertop2tuned --help
```

3.2. PERFORMANCE TUNING WITH TUNED AND TUNED-ADM

The **tuned** tuning service can adapt the operating system to perform better under certain workloads by setting a tuning profile. The **tuned-adm** command-line tool allows users to switch between different tuning profiles.

tuned Profiles Overview

Several pre-defined profiles are included for common use cases, but **tuned** also enables you to define custom profiles, which can be either based on one of the pre-defined profiles, or defined from scratch. In Red Hat Enterprise Linux 7, the default profile is **throughput-performance**.

The profiles provided with **tuned** are divided into two categories: power-saving profiles, and performance-boosting profiles. The performance-boosting profiles include profiles focus on the following aspects:

- low latency for storage and network
- high throughput for storage and network
- virtual machine performance

- virtualization host performance

tuned Boot Loader plug-in

You can use the **`tuned Bootloader plug-in`** to add parameters to the kernel (boot or dracut) command line. Note that only the GRUB 2 boot loader is supported and a reboot is required to apply profile changes. For example, to add the ***`quiet`*** parameter to a **tuned** profile, include the following lines in the **`tuned.conf`** file:

```
[bootloader]
cmdline=quiet
```

Switching to another profile or manually stopping the **tuned** service removes the additional parameters. If you shut down or reboot the system, the kernel parameters persist in the **`grub.cfg`** file.

Environment Variables and Expanding tuned Built-In Functions

If you run **`tuned-adm profile`** ***`profile_name`*** and then **`grub2-mkconfig -o`** ***`profile_path`*** after updating GRUB 2 configuration, you can use Bash environment variables, which are expanded after running **`grub2-mkconfig`**. For example, the following environment variable is expanded to **`nfsroot=/root`**:

```
[bootloader]
cmdline="nfsroot=$HOME"
```

You can use **`tuned`** variables as an alternative to environment variables. In the following example, **`${isolated_cores}`** expands to **`1,2`**, so the kernel boots with the **`isolcpus=1,2`** parameter:

```
[variables]
isolated_cores=1,2

[bootloader]
cmdline=isolcpus=${isolated_cores}
```

In the following example, **`${non_isolated_cores}`** expands to **`0,3-5`**, and the **`cpulist_invert`** built-in function is called with the **`0,3-5`** arguments:

```
[variables]
non_isolated_cores=0,3-5

[bootloader]
cmdline=isolcpus=${f:cpulist_invert:${non_isolated_cores}}
```

The **`cpulist_invert`** function inverts the list of CPUs. For a 6-CPU machine, the inversion is**`1,2`**, and the kernel boots with the **`isolcpus=1,2`** command-line parameter.

Using **tuned** environment variables reduces the amount of necessary typing. You can also use various built-in functions together with **tuned** variables. If the built-in functions do not satisfy your needs, you can create custom functions in Python and add them to **tuned** in the form of plug-ins. Variables and built-in functions are expanded at run time when the tuned profile is activated.

The variables can be specified in a separate file. You can, for example, add the following lines to **`tuned.conf`**:

```
[variables]
include=/etc/tuned/my-variables.conf
```

```
[bootloader]
cmdline=isolcpus=${isolated_cores}
```

If you add **isolated_cores=1,2** to the **/etc/tuned/*my-variables.conf*** file, the kernel boots with the **isolcpus=1,2** parameter.

Modifying Default System tuned Profiles

There are two ways of modifying the default system **tuned** profiles. You can either create a new **tuned** profile directory, or copy the directory of a system profile and edit the profile as needed.

Procedure 3.1. Creating a New Tuned Profile Directory

1. In **/etc/tuned/**, create a new directory named the same as the profile you want to create: **/etc/tuned/*my_profile_name*/**.

2. In the new directory, create a file named **tuned.conf**, and include the following lines at the top:

   ```
   [main]
   include=profile_name
   ```

3. Include your profile modifications. For example, to use the settings from the **throughput-performance** profile with the value of **vm.swappiness** set to 5, instead of default 10, include the following lines:

   ```
   [main]
   include=throughput-performance

   [sysctl]
   vm.swappiness=5
   ```

4. To activate the profile, run:

   ```
   # tuned-adm profile my_profile_name
   ```

Creating a directory with a new **tuned.conf** file enables you to keep all your profile modifications after system **tuned** profiles are updated.

Alternatively, copy the directory with a system profile from **/user/lib/tuned/** to **/etc/tuned/**. For example:

```
# cp -r /usr/lib/tuned/throughput-performance /etc/tuned
```

Then, edit the profile in **/etc/tuned** according to your needs. Note that if there are two profiles of the same name, the profile located in **/etc/tuned/** is loaded. The disadvantage of this approach is that if a system profile is updated after a **tuned** upgrade, the changes will not be reflected in the now-outdated modified version.

Resources

For more information, see Section A.4, “tuned” and Section A.5, “tuned-adm”. For detailed information on using **tuned** and **tuned-adm**, see the tuned(8) and tuned-adm(1) manual pages.

CHAPTER 4. TUNA

You can use the **Tuna** tool to adjust scheduler tunables, tune thread priority, IRQ handlers, and isolate CPU cores and sockets. Tuna aims to reduce the complexity of performing tuning tasks.

After installing the tuna package, use the **`tuna`** command without any arguments to start the Tuna graphical user interface (GUI). Use the **`tuna -h`** command to display available command-line interface (CLI) options. Note that the tuna(8) manual page distinguishes between action and modifier options.

The Tuna GUI and CLI provide equivalent functionality. The GUI displays the CPU topology on one screen to help you identify problems. The Tuna GUI also allows you to make changes to the running threads, and see the results of those changes immediately. In the CLI, Tuna accepts multiple command-line parameters and processes them sequentially. You can use such commands in application initialization scripts as configuration commands.

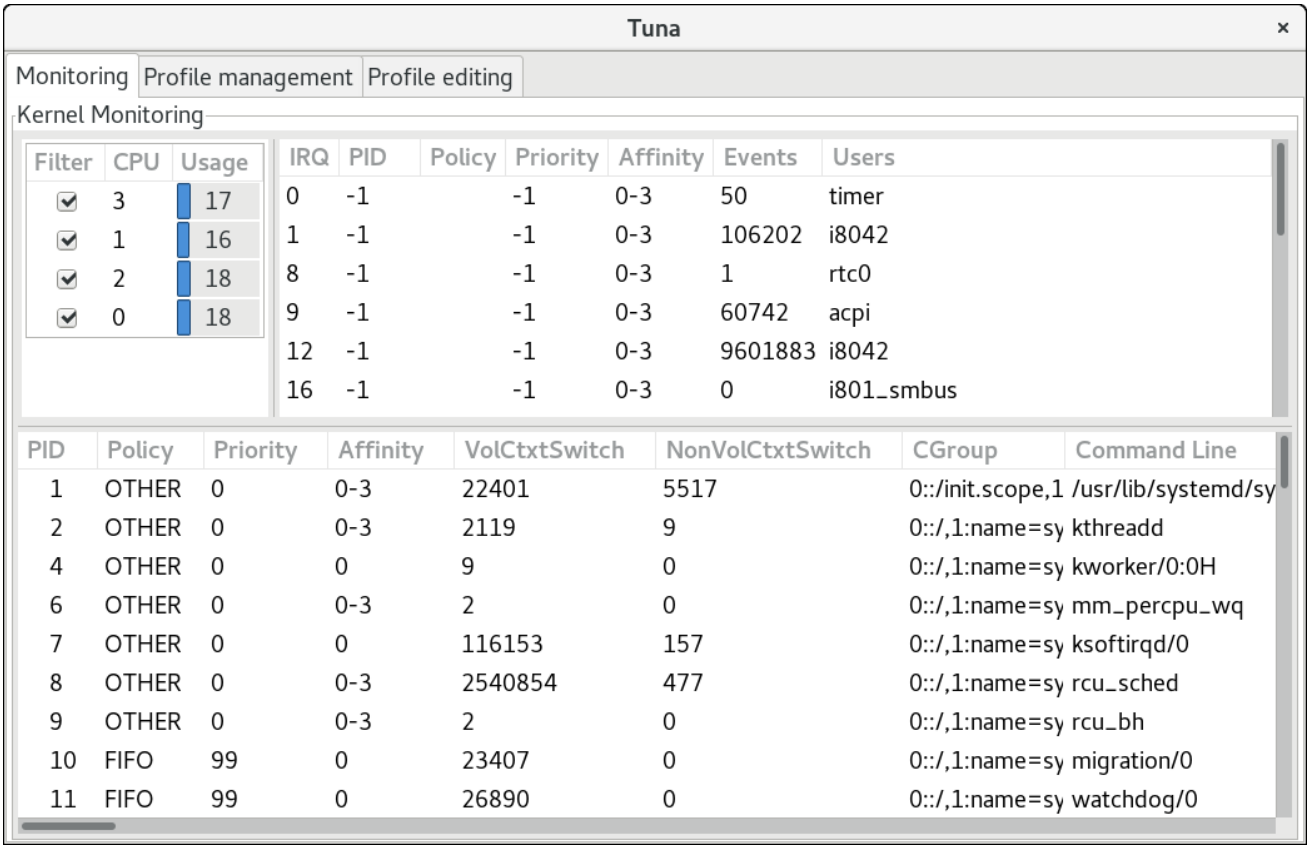

The Monitoring tab of the Tuna GUI

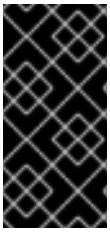

IMPORTANT

Use the **`tuna --save=`*`filename`*** command with a descriptive file name to save the current configuration. Note that this command does *not* save every option that Tuna can change, but saves the kernel thread changes only. Any processes that are not currently running when they are changed are not saved.

4.1. REVIEWING THE SYSTEM WITH TUNA

Before you make any changes, you can use Tuna to show you what is currently happening on the system.

To view the current policies and priorities, use the **`tuna --show_threads`** command:

```
# tuna --show_threads
      thread
pid   SCHED_ rtpri affinity            cmd
1      OTHER     0      0,1           init
2       FIFO    99        0     migration/0
3      OTHER     0        0     ksoftirqd/0
4       FIFO    99        0      watchdog/0
```

To show only a specific thread corresponding to a PID or matching a command name, add the **--threads** option before **--show_threads**:

```
# tuna --threads=pid_or_cmd_list --show_threads
```

The *pid_or_cmd_list* argument is a list of comma-separated PIDs or command-name patterns.

To view the current interrupt requests (IRQs) and their affinity, use the **tuna --show_irqs** command:

```
# tuna --show_irqs
# users            affinity
0 timer                  0
1 i8042                  0
7 parport0               0
```

To show only a specific interrupt request corresponding to an IRQ number or matching an IRQ user name, add the **--irqs** option before **--show_irqs**:

```
# tuna --irqs=number_or_user_list --show_irqs
```

The *number_or_user_list* argument is a list of comma-separated IRQ numbers or user-name patterns.

4.2. TUNING CPUS WITH TUNA

Tuna commands can target individual CPUs. To list the CPUs on your system, see the Monitoring tab in Tuna GUI or the **/proc/cpuinfo** file for detailed information.

To specify the list of CPUs to be affected by your command, use:

```
# tuna --cpus=cpu_list --command
```

Isolating a CPU causes all tasks currently running on that CPU to move to the next available CPU. To isolate a CPU, use:

```
# tuna --cpus=cpu_list --isolate
```

Including a CPU allows threads to run on the specified CPU. To include a CPU, use:

```
# tuna --cpus=cpu_list --include
```

The *cpu_list* argument is a list of comma-separated CPU numbers. For example, **--cpus=0,2**.

4.3. TUNING IRQS WITH TUNA

To see the list of IRQs currently running on your system, see the Monitoring tab in Tuna GUI or the **/proc/interrpupts** file. You can also use the **tuna --show_irqs** command.

To specify the list of IRQs to be affected by your command, use the **--irqs** parameter:

```
# tuna --irqs=irq_list --command
```

To move an interrupt to a specified CPU, use the **--move** parameter:

```
# tuna --irqs=irq_list --cpus=cpu_list --move
```

The *irq_list* argument is a list of comma-separated IRQ numbers or user-name patterns.

The *cpu_list* argument is a list of comma-separated CPU numbers. For example, **--cpus=0,2**.

For example, to target all interrupts whose names start with **sfc1** and spread them over two CPUs:

```
# tuna --irqs=sfc1\* --cpus=7,8 --move --spread
```

To verify the changes you set, use the **--show_irqs** parameter both before and after modifying the IRQs with the **--move** parameter:

```
# tuna --irqs=128 --show_irqs

   # users           affinity
 128 iwlwifi          0,1,2,3

# tuna --irqs=128 --cpus=3 --move

# tuna --irqs=128 --show_irqs

   # users           affinity
 128 iwlwifi                3
```

This allows you to compare the state of the selected IRQs before and after your changes.

NOTE

Using the Tuna GUI might be more convenient in certain scenarios. Moving IRQs and threads by specifying the CPUs they are to run on can be time-consuming and difficult because it involves multiple steps of creating CPU masks. The Tuna GUI automates the process. In the Tuna GUI, you can also select threads and IRQs and drag them over the intended CPUs, which can make changing the topology easier.

4.4. TUNING TASKS WITH TUNA

To change policy and priority information on threads, use the **--priority** parameter:

```
# tuna --threads=pid_or_cmd_list --priority=[policy:]rt_priority
```

- The *pid_or_cmd_list* argument is a list of comma-separated PIDs or command-name patterns.
- Set the *policy* to **RR** for round-robin, **FIFO** for first in, first out, or **OTHER** for the default policy.

For an overview of the scheduling policies, see Section 5.3.6, “Tuning Scheduling Policy”.

- Set the *rt_priority* in the range 1–99. 1 is the lowest priority, and 99 is the highest priority.

For example:

```
# tuna --threads=7861 --priority=RR:40
```

To verify the changes you set, use the **--show_threads** parameter both before and after the modifying **--priority** parameter:

```
# tuna --threads=sshd --show_threads --priority=RR:40 --show_threads

                    thread       ctxt_switches
  pid SCHED_ rtpri affinity voluntary nonvoluntary             cmd
 1034   OTHER     0  0,1,2,3        12           17            sshd
                    thread       ctxt_switches
  pid SCHED_ rtpri affinity voluntary nonvoluntary             cmd
 1034      RR    40  0,1,2,3        12           17            sshd
```

This allows you to compare the state of the selected threads before and after your changes.

4.5. EXAMPLES OF USING TUNA

Example 4.1. Assigning Tasks to Specific CPUs

The following example uses a system with four or more processors, and shows how to make all **ssh** threads run on CPU 0 and 1, and all **http** threads on CPU 2 and 3.

```
# tuna --cpus=0,1 --threads=ssh\* --move --cpus=2,3 --threads=http\* --
move
```

The above example command performs the following operations sequentially:

1. Select CPUs 0 and 1.
2. Select all threads that begin with **ssh**.
3. Move the selected threads to the selected CPUs. Tuna sets the affinity mask of threads starting with **ssh** to the appropriate CPUs. The CPUs can be expressed numerically as 0 and 1, in hex mask as 0x3, or in binary as 11.
4. Reset the CPU list to 2 and 3.
5. Select all threads that begin with **http**.
6. Move the selected threads to the selected CPUs. Tuna sets the affinity mask of threads starting with **http** to the appropriate CPUs. The CPUs can be expressed numerically as 2 and 3, in hex mask as 0xC, or in binary as 1100.

Example 4.2. Viewing Current Configurations

The following example uses the **--show_threads** (**-P**) parameter to display the current configuration, and then tests if the requested changes were made as expected.

```
# tuna --threads=gnome-sc\* \
       --show_threads \
       --cpus=0 \
       --move \
       --show_threads \
       --cpus=1 \
       --move \
       --show_threads \
       --cpus=+0 \
       --move \
       --show_threads

                     thread       ctxt_switches
  pid SCHED_ rtpri affinity voluntary nonvoluntary             cmd
 3861   OTHER     0      0,1     33997           58 gnome-screensav
                     thread       ctxt_switches
  pid SCHED_ rtpri affinity voluntary nonvoluntary             cmd
 3861   OTHER     0        0     33997           58 gnome-screensav
                     thread       ctxt_switches
  pid SCHED_ rtpri affinity voluntary nonvoluntary             cmd
 3861   OTHER     0        1     33997           58 gnome-screensav
                     thread       ctxt_switches
  pid SCHED_ rtpri affinity voluntary nonvoluntary             cmd
 3861   OTHER     0      0,1     33997           58 gnome-screensav
```

The above example command performs the following operations sequentially:

1. Select all threads that begin with **gnome-sc**.
2. Show the selected threads to enable the user to verify their affinity mask and RT priority.
3. Select CPU 0.
4. Move the **gnome-sc** threads to the selected CPU (CPU 0).
5. Show the result of the move.
6. Reset the CPU list to CPU 1.
7. Move the **gnome-sc** threads to the selected CPU (CPU 1).
8. Show the result of the move.
9. Add CPU 0 to the CPU list.
10. Move the **gnome-sc** threads to the selected CPUs (CPUs 0 and 1).
11. Show the result of the move.

CHAPTER 5. CPU

This chapter outlines CPU hardware details and configuration options that affect application performance in Red Hat Enterprise Linux 7. Section 5.1, “Considerations” discusses the CPU related factors that affect performance. Section 5.2, “Monitoring and Diagnosing Performance Problems” teaches you how to use Red Hat Enterprise Linux 7 tools to diagnose performance problems related to CPU hardware or configuration details. Section 5.3, “Configuration Suggestions” discusses the tools and strategies you can use to solve CPU related performance problems in Red Hat Enterprise Linux 7.

5.1. CONSIDERATIONS

Read this section to gain an understanding of how system and application performance is affected by the following factors:

- How processors are connected to each other and to related resources like memory.
- How processors schedule threads for execution.
- How processors handle interrupts in Red Hat Enterprise Linux 7.

5.1.1. System Topology

In modern computing, the idea of a *central* processing unit is a misleading one, as most modern systems have multiple processors. How these processors are connected to each other and to other system resources — the *topology* of the system — can greatly affect system and application performance, and the tuning considerations for a system.

There are two primary types of topology used in modern computing:

Symmetric Multi-Processor (SMP) topology

SMP topology allows all processors to access memory in the same amount of time. However, because shared and equal memory access inherently forces serialized memory accesses from all the CPUs, SMP system scaling constraints are now generally viewed as unacceptable. For this reason, practically all modern server systems are NUMA machines.

Non-Uniform Memory Access (NUMA) topology

NUMA topology was developed more recently than SMP topology. In a NUMA system, multiple processors are physically grouped on a socket. Each socket has a dedicated area of memory, and processors that have local access to that memory are referred to collectively as a node.

Processors on the same node have high speed access to that node's memory bank, and slower access to memory banks not on their node. Therefore, there is a performance penalty to accessing non-local memory.

Given this performance penalty, performance sensitive applications on a system with NUMA topology should access memory that is on the same node as the processor executing the application, and should avoid accessing remote memory wherever possible.

When tuning application performance on a system with NUMA topology, it is therefore important to consider where the application is being executed, and which memory bank is closest to the point of execution.

In a system with NUMA topology, the **/sys** file system contains information about how processors, memory, and peripheral devices are connected. The **/sys/devices/system/cpu** directory

contains details about how processors in the system are connected to each other. The **/sys/devices/system/node** directory contains information about NUMA nodes in the system, and the relative distances between those nodes.

5.1.1.1. Determining System Topology

There are a number of commands that can help you understand the topology of your system. The **numactl --hardware** command gives an overview of your system's topology.

```
$ numactl --hardware
available: 4 nodes (0-3)
node 0 cpus: 0 4 8 12 16 20 24 28 32 36
node 0 size: 65415 MB
node 0 free: 43971 MB
node 1 cpus: 2 6 10 14 18 22 26 30 34 38
node 1 size: 65536 MB
node 1 free: 44321 MB
node 2 cpus: 1 5 9 13 17 21 25 29 33 37
node 2 size: 65536 MB
node 2 free: 44304 MB
node 3 cpus: 3 7 11 15 19 23 27 31 35 39
node 3 size: 65536 MB
node 3 free: 44329 MB
node distances:
node   0   1   2   3
  0:  10  21  21  21
  1:  21  10  21  21
  2:  21  21  10  21
  3:  21  21  21  10
```

The **lscpu** command, provided by the util-linux package, gathers information about the CPU architecture, such as the number of CPUs, threads, cores, sockets, and NUMA nodes.

```
$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
CPU(s):                40
On-line CPU(s) list:   0-39
Thread(s) per core:    1
Core(s) per socket:    10
Socket(s):             4
NUMA node(s):          4
Vendor ID:             GenuineIntel
CPU family:            6
Model:                 47
Model name:            Intel(R) Xeon(R) CPU E7- 4870  @ 2.40GHz
Stepping:              2
CPU MHz:               2394.204
BogoMIPS:              4787.85
Virtualization:        VT-x
L1d cache:             32K
L1i cache:             32K
L2 cache:              256K
```

```
L3 cache:              30720K
NUMA node0 CPU(s):     0,4,8,12,16,20,24,28,32,36
NUMA node1 CPU(s):     2,6,10,14,18,22,26,30,34,38
NUMA node2 CPU(s):     1,5,9,13,17,21,25,29,33,37
NUMA node3 CPU(s):     3,7,11,15,19,23,27,31,35,39
```

The **lstopo** command, provided by the hwloc package, creates a graphical representation of your system. The **lstopo-no-graphics** command provides detailed textual output.

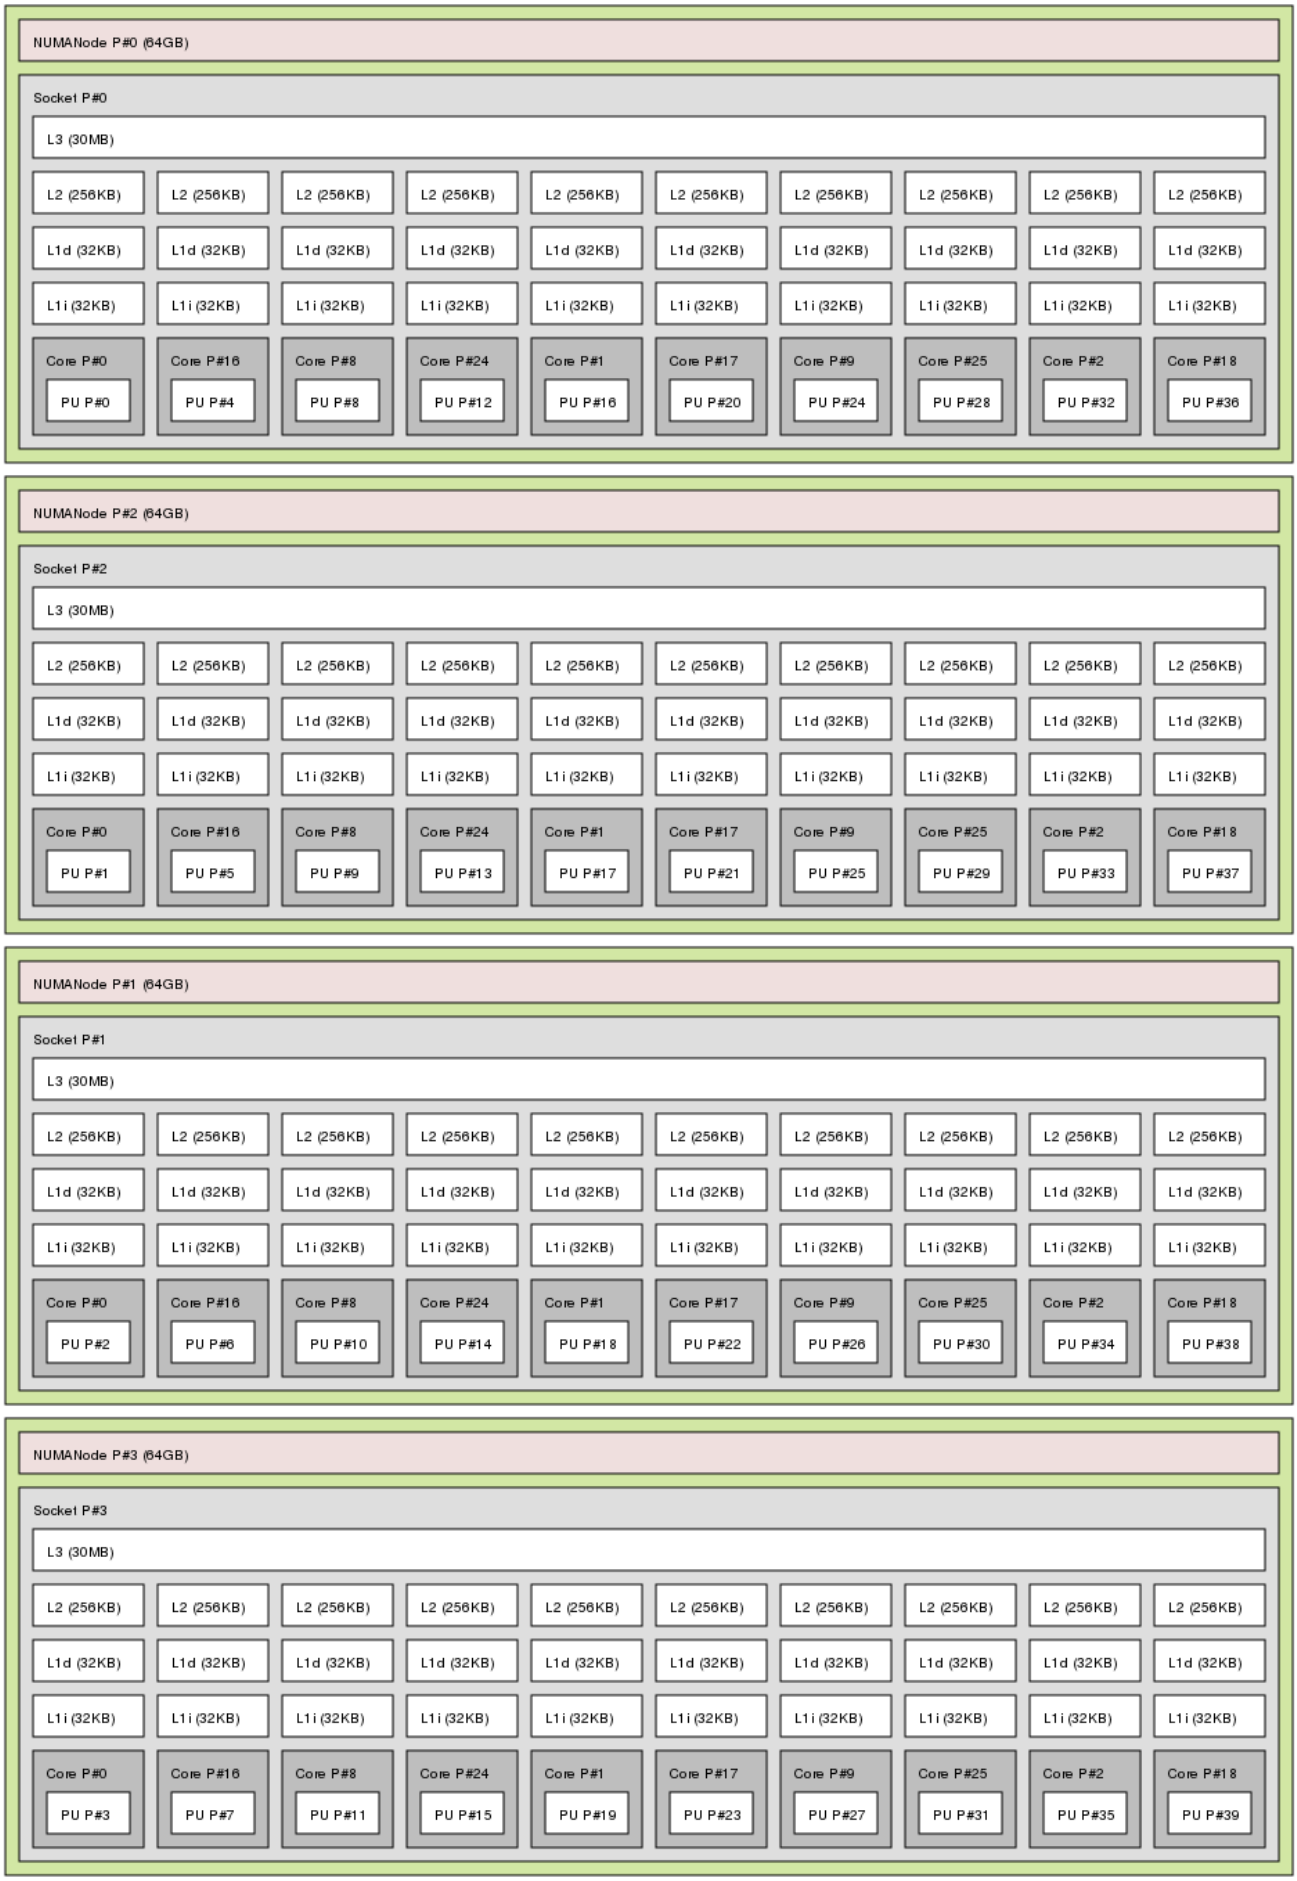

Output of lstopo command

5.1.2. Scheduling

In Red Hat Enterprise Linux, the smallest unit of process execution is called a *thread*. The system scheduler determines which processor runs a thread, and for how long the thread runs. However, because the scheduler's primary concern is to keep the system busy, it may not schedule threads optimally for application performance.

For example, say an application on a NUMA system is running on Node A when a processor on Node B becomes available. To keep the processor on Node B busy, the scheduler moves one of the application's threads to Node B. However, the application thread still requires access to memory on Node A. Because the thread is now running on Node B, and Node A memory is no longer local to the thread, it will take longer to access. It may take longer for the thread to finish running on Node B than it would have taken to wait for a processor on Node A to become available, and to execute the thread on the original node with local memory access.

Performance sensitive applications often benefit from the designer or administrator determining where threads are run. For details about how to ensure threads are scheduled appropriately for the needs of performance sensitive applications, see Section 5.3.6, “Tuning Scheduling Policy”.

5.1.2.1. Kernel Ticks

In previous versions of Red Hat Enterprise Linux, the Linux kernel interrupted each CPU on a regular basis to check what work needed to be done. It used the results to make decisions about process scheduling and load balancing. This regular interruption was known as a kernel *tick*.

This tick occurred regardless of whether there was work for the core to do. This meant that even idle cores were forced into higher power states on a regular basis (up to 1000 times per second) to respond to the interrupts. This prevented the system from effectively using deep sleep states included in recent generations of x86 processors.

In Red Hat Enterprise Linux 6 and 7, by default, the kernel no longer interrupts idle CPUs, which tend to be in low power states. This behavior is known as the tickless kernel. Where one or fewer tasks are running, periodic interrupts have been replaced with on-demand interrupts, allowing CPUs to remain in an idle or low power state for longer, and reducing power usage.

Red Hat Enterprise Linux 7 offers a dynamic tickless option (**`nohz_full`**) to further improve determinism by reducing kernel interference with user-space tasks. This option can be enabled on specified cores with the ***`nohz_full`*** kernel parameter. When this option is enabled on a core, all timekeeping activities are moved to non-latency-sensitive cores. This can be useful for high performance computing and realtime computing workloads where user-space tasks are particularly sensitive to microsecond-level latencies associated with the kernel timer tick.

For details on how to enable the dynamic tickless behavior in Red Hat Enterprise Linux 7, see Section 5.3.1, “Configuring Kernel Tick Time”.

5.1.3. Interrupt Request (IRQ) Handling

An interrupt request or IRQ is a signal for immediate attention sent from a piece of hardware to a processor. Each device in a system is assigned one or more IRQ numbers to allow it to send unique interrupts. When interrupts are enabled, a processor that receives an interrupt request will immediately pause execution of the current application thread in order to address the interrupt request.

Because they halt normal operation, high interrupt rates can severely degrade system performance. It is possible to reduce the amount of time taken by interrupts by configuring interrupt affinity or by sending a number of lower priority interrupts in a batch (*coalescing* a number of interrupts).

For more information about tuning interrupt requests, see Section 5.3.7, “Setting Interrupt Affinity” or Section 5.3.8, “Configuring CPU, Thread, and Interrupt Affinity with Tuna”. For information specific to network interrupts, see Chapter 8, *Networking*.

5.2. MONITORING AND DIAGNOSING PERFORMANCE PROBLEMS

Red Hat Enterprise Linux 7 provides a number of tools that are useful for monitoring system performance and diagnosing performance problems related to processors and their configuration. This section outlines the available tools and gives examples of how to use them to monitor and diagnose processor related performance issues.

5.2.1. turbostat

Turbostat prints counter results at specified intervals to help administrators identify unexpected behavior in servers, such as excessive power usage, failure to enter deep sleep states, or system management interrupts (SMIs) being created unnecessarily.

The **turbostat** tool is part of the kernel-tools package. It is supported for use on systems with AMD64 and Intel® 64 processors. It requires root privileges to run, and processor support for invariant time stamp counters, and APERF and MPERF model specific registers.

For usage examples, see the man page:

```
$ man turbostat
```

5.2.2. numastat

IMPORTANT

This tool received substantial updates in the Red Hat Enterprise Linux 6 life cycle. While the default output remains compatible with the original tool written by Andi Kleen, supplying any options or parameters to numastat significantly changes the format of its output.

The **numastat** tool displays per-NUMA node memory statistics for processes and the operating system and shows administrators whether process memory is spread throughout a system or centralized on specific nodes.

Cross reference **numastat** output with per-processor **top** output to confirm that process threads are running on the same node from which process memory is allocated.

Numastat is provided by the numactl package. For further information about **numastat** output, see the man page:

```
$ man numastat
```

5.2.3. /proc/interrupts

The **`/proc/interrupts`** file lists the number of interrupts sent to each processor from a particular I/O device. It displays the interrupt request (IRQ) number, the number of that type of interrupt request handled by each processor in the system, the type of interrupt sent, and a comma-separated list of devices that respond to the listed interrupt request.

If a particular application or device is generating a large number of interrupt requests to be handled by a remote processor, its performance is likely to suffer. In this case, poor performance can be alleviated by having a processor on the same node as the application or device handle the interrupt requests. For details on how to assign interrupt handling to a specific processor, see Section 5.3.7, “Setting Interrupt Affinity”.

5.3. CONFIGURATION SUGGESTIONS

Red Hat Enterprise Linux provides a number of tools to assist administrators in configuring the system. This section outlines the available tools and provides examples of how they can be used to solve processor related performance problems in Red Hat Enterprise Linux 7.

5.3.1. Configuring Kernel Tick Time

By default, Red Hat Enterprise Linux 7 uses a tickless kernel, which does not interrupt idle CPUs in order to reduce power usage and allow newer processors to take advantage of deep sleep states.

Red Hat Enterprise Linux 7 also offers a dynamic tickless option (disabled by default), which is useful for very latency-sensitive workloads, such as high performance computing or realtime computing.

To enable dynamic tickless behavior in certain cores, specify those cores on the kernel command line with the ***nohz_full*** parameter. On a 16 core system, specifying **nohz_full=1-15** enables dynamic tickless behavior on cores 1 through 15, moving all timekeeping to the only unspecified core (core 0). This behavior can be enabled either temporarily at boot time, or persistently in the **/etc/default/grub** file. For persistent behavior, run the **grub2-mkconfig -o /boot/grub2/grub.cfg** command to save your configuration.

Enabling dynamic tickless behavior does require some manual administration.

- When the system boots, you must manually move rcu threads to the non-latency-sensitive core, in this case core 0.

  ```
  # for i in `pgrep rcu[^c]` ; do taskset -pc 0 $i ; done
  ```

- Use the ***isolcpus*** parameter on the kernel command line to isolate certain cores from user-space tasks.

- Optionally, set CPU affinity for the kernel's write-back bdi-flush threads to the housekeeping core:

  ```
  echo 1 > /sys/bus/workqueue/devices/writeback/cpumask
  ```

Verify that the dynamic tickless configuration is working correctly by executing the following command, where *stress* is a program that spins on the CPU for 1 second.

```
# perf stat -C 1 -e irq_vectors:local_timer_entry taskset -c 1 stress -t 1
-c 1
```

One possible replacement for *stress* is a script that runs something like **while :; do d=1; done**.

The default kernel timer configuration shows 1000 ticks on a busy CPU:

```
# perf stat -C 1 -e irq_vectors:local_timer_entry taskset -c 1 stress -t 1
-c 1
```

```
1000 irq_vectors:local_timer_entry
```

With the dynamic tickless kernel configured, you should see 1 tick instead:

```
# perf stat -C 1 -e irq_vectors:local_timer_entry taskset -c 1 stress -t 1
-c 1
1 irq_vectors:local_timer_entry
```

5.3.2. Setting Hardware Performance Policy (x86_energy_perf_policy)

The **x86_energy_perf_policy** tool allows administrators to define the relative importance of performance and energy efficiency. This information can then be used to influence processors that support this feature when they select options that trade off between performance and energy efficiency.

By default, it operates on all processors in **`performance`** mode. It requires processor support, which is indicated by the presence of **`CPUID.06H.ECX.bit3`**, and must be run with root privileges.

x86_energy_perf_policy is provided by the kernel-tools package. For details of how to use **x86_energy_perf_policy**, see Section A.9, “x86_energy_perf_policy” or refer to the man page:

```
$ man x86_energy_perf_policy
```

5.3.3. Setting Process Affinity with taskset

The **taskset** tool is provided by the util-linux package. **Taskset** allows administrators to retrieve and set the processor affinity of a running process, or launch a process with a specified processor affinity.

IMPORTANT

taskset does not guarantee local memory allocation. If you require the additional performance benefits of local memory allocation, Red Hat recommends using **numactl** instead of **taskset**.

For more information about **taskset**, see Section A.15, “taskset” or the man page:

```
$ man taskset
```

5.3.4. Managing NUMA Affinity with numactl

Administrators can use **numactl** to run a process with a specified scheduling or memory placement policy. **Numactl** can also set a persistent policy for shared memory segments or files, and set the processor affinity and memory affinity of a process.

In a system with NUMA topology, a processor's memory access slows as the distance between the processor and the memory bank increases. Therefore, it is important to configure applications that are sensitive to performance so that they allocate memory from the closest possible memory bank. It is best to use memory and CPUs that are in the same NUMA node.

Multi-threaded applications that are sensitive to performance may benefit from being configured to execute on a specific NUMA node rather than a specific processor. Whether this is suitable depends on your system and the requirements of your application. If multiple application threads access the same cached data, then configuring those threads to execute on the same processor may be suitable.

However, if multiple threads that access and cache different data execute on the same processor, each thread may evict cached data accessed by a previous thread. This means that each thread 'misses' the cache, and wastes execution time fetching data from memory and replacing it in the cache. You can use the **perf** tool, as documented in Section A.6, “perf”, to check for an excessive number of cache misses.

Numactl provides a number of options to assist you in managing processor and memory affinity. See Section A.11, “numastat” or the man page for details:

```
$ man numactl
```

NOTE

The **numactl** package includes the **libnuma** library. This library offers a simple programming interface to the NUMA policy supported by the kernel, and can be used for more fine-grained tuning than the **numactl** application. For more information, see the man page:

```
$ man numa
```

5.3.5. Automatic NUMA Affinity Management with numad

numad is an automatic NUMA affinity management daemon. It monitors NUMA topology and resource usage within a system in order to dynamically improve NUMA resource allocation and management.

numad also provides a pre-placement advice service that can be queried by various job management systems to provide assistance with the initial binding of CPU and memory resources for their processes. This pre-placement advice is available regardless of whether numad is running as an executable or a service.

For details of how to use **numad**, see Section A.13, “numad” or refer to the man page:

```
$ man numad
```

5.3.6. Tuning Scheduling Policy

The Linux scheduler implements a number of scheduling policies, which determine where and for how long a thread runs. There are two major categories of scheduling policies: normal policies and realtime policies. Normal threads are used for tasks of normal priority. Realtime policies are used for time-sensitive tasks that must complete without interruptions.

Realtime threads are not subject to time slicing. This means they will run until they block, exit, voluntarily yield, or are pre-empted by a higher priority thread. The lowest priority realtime thread is scheduled before any thread with a normal policy.

5.3.6.1. Scheduling Policies

5.3.6.1.1. Static Priority Scheduling with SCHED_FIFO

SCHED_FIFO (also called static priority scheduling) is a realtime policy that defines a fixed priority for each thread. This policy allows administrators to improve event response time and reduce latency, and is recommended for time sensitive tasks that do not run for an extended period of time.

When **SCHED_FIFO** is in use, the scheduler scans the list of all **SCHED_FIFO** threads in priority order and schedules the highest priority thread that is ready to run. The priority level of a **SCHED_FIFO** thread can be any integer from 1 to 99, with 99 treated as the highest priority. Red Hat recommends starting at a low number and increasing priority only when you identify latency issues.

WARNING

Because realtime threads are not subject to time slicing, Red Hat does not recommend setting a priority of 99. This places your process at the same priority level as migration and watchdog threads; if your thread goes into a computational loop and these threads are blocked, they will not be able to run. Systems with a single processor will eventually hang in this situation.

Administrators can limit **SCHED_FIFO** bandwidth to prevent realtime application programmers from initiating realtime tasks that monopolize the processor.

/proc/sys/kernel/sched_rt_period_us

This parameter defines the time period in microseconds that is considered to be one hundred percent of processor bandwidth. The default value is **1000000** µs, or 1 second.

/proc/sys/kernel/sched_rt_runtime_us

This parameter defines the time period in microseconds that is devoted to running realtime threads. The default value is **950000** µs, or 0.95 seconds.

5.3.6.1.2. Round Robin Priority Scheduling with SCHED_RR

SCHED_RR is a round-robin variant of **SCHED_FIFO**. This policy is useful when multiple threads need to run at the same priority level.

Like **SCHED_FIFO**, **SCHED_RR** is a realtime policy that defines a fixed priority for each thread. The scheduler scans the list of all **SCHED_RR** threads in priority order and schedules the highest priority thread that is ready to run. However, unlike **SCHED_FIFO**, threads that have the same priority are scheduled round-robin style within a certain time slice.

You can set the value of this time slice in milliseconds with the ***sched_rr_timeslice_ms*** kernel parameter (**/proc/sys/kernel/sched_rr_timeslice_ms**). The lowest value is 1 millisecond.

5.3.6.1.3. Normal Scheduling with SCHED_OTHER

SCHED_OTHER is the default scheduling policy in Red Hat Enterprise Linux 7. This policy uses the Completely Fair Scheduler (CFS) to allow fair processor access to all threads scheduled with this policy. This policy is most useful when there are a large number of threads or data throughput is a priority, as it allows more efficient scheduling of threads over time.

When this policy is in use, the scheduler creates a dynamic priority list based partly on the niceness value of each process thread. Administrators can change the niceness value of a process, but cannot change the scheduler's dynamic priority list directly.

For details about changing process niceness, see the *Red Hat Enterprise Linux 7 Deployment Guide*, available from http://access.redhat.com/site/documentation/Red_Hat_Enterprise_Linux/.

5.3.6.2. Isolating CPUs

You can isolate one or more CPUs from the scheduler with the ***isolcpus*** boot parameter. This prevents the scheduler from scheduling any user-space threads on this CPU.

Once a CPU is isolated, you must manually assign processes to the isolated CPU, either with the CPU affinity system calls or the numactl command.

To isolate the third and sixth to eighth CPUs on your system, add the following to the kernel command line:

```
isolcpus=2,5-7
```

You can also use the **Tuna** tool to isolate a CPU. **Tuna** can isolate a CPU at any time, not just at boot time. However, this method of isolation is subtly different from the ***isolcpus*** parameter, and does not currently achieve the performance gains associated with ***isolcpus***. See Section 5.3.8, “Configuring CPU, Thread, and Interrupt Affinity with Tuna” for more details about this tool.

5.3.7. Setting Interrupt Affinity

Interrupt requests have an associated affinity property, ***smp_affinity***, that defines the processors that will handle the interrupt request. To improve application performance, assign interrupt affinity and process affinity to the same processor, or processors on the same core. This allows the specified interrupt and application threads to share cache lines.

The interrupt affinity value for a particular interrupt request is stored in the associated **/proc/irq/*irq_number*/smp_affinity** file. ***smp_affinity*** is stored as a hexadecimal bit mask representing all processors in the system. The default value is **f**, meaning that an interrupt request can be handled on any processor in the system. Setting this value to **1** means that only processor 0 can handle the interrupt.

On systems with more than 32 processors, you must delimit ***smp_affinity*** values for discrete 32 bit groups. For example, if you want only the first 32 processors of a 64 processor system to service an interrupt request, you could run:

```
# echo 0xffffffff,00000000 > /proc/irq/IRQ_NUMBER/smp_affinity
```

Alternatively, if the BIOS exports its NUMA topology, the **irqbalance** service can use that information to serve interrupt requests on the node that is local to the hardware requesting service. For details about **irqbalance**, see Section A.1, “irqbalance”

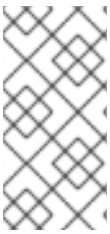

NOTE

On systems that support interrupt steering, modifying the ***smp_affinity*** of an interrupt request sets up the hardware so that the decision to service an interrupt with a particular processor is made at the hardware level with no intervention from the kernel. For more information about interrupt steering, see Chapter 8, *Networking*.

5.3.8. Configuring CPU, Thread, and Interrupt Affinity with Tuna

Tuna is a tool for tuning running processes and can control CPU, thread, and interrupt affinity, and also provides a number of actions for each type of entity it can control. For information about **Tuna**, see Chapter 4, *Tuna*.

CHAPTER 6. MEMORY

This chapter outlines the memory management capabilities of Red Hat Enterprise Linux 7. Section 6.1, “Considerations” discusses memory related factors that affect performance. Section 6.2, “Monitoring and Diagnosing Performance Problems” teaches you how to use Red Hat Enterprise Linux 7 tools to diagnose performance problems related to memory utilization or configuration details. Section 6.5, “Configuring System Memory Capacity” and Section 6.3, “Configuring HugeTLB Huge Pages” discuss the tools and strategies you can use to solve memory related performance problems in Red Hat Enterprise Linux 7.

6.1. CONSIDERATIONS

By default, Red Hat Enterprise Linux 7 is optimized for moderate workloads. If your application or workload requires a large amount of memory, changing the way your system manages virtual memory may improve the performance of your application.

6.1.1. Larger Page Size

Physical memory is managed in chunks called pages. On most architectures supported by Red Hat Enterprise Linux 7, the default size of a memory page is 4 KB. This default page size has proved to be suitable for general-purpose operating systems, such as Red Hat Enterprise Linux 7, which support many different kinds of workloads.

However, specific applications can benefit from using larger page sizes in certain cases. For example, an application that works with a large and relatively fixed data set of hundreds of megabytes or even dozens of gigabytes can have performance issues when using 4 KB pages. Such data sets can require hundreds of thousands of 4 KB pages, which can lead to overhead in the operating system and the CPU.

Red Hat Enterprise Linux 7 enables the use of larger page sizes for applications working with big data sets. Using larger page sizes can improve the performance of such applications.

Two different large page features are available in Red Hat Enterprise Linux 7: the **`HugeTLB`** feature, also called **`static huge pages`** in this guide, and the **`Transparent Huge Page`** feature.

6.1.2. Translation Lookaside Buffer Size

Reading address mappings from the page table is time-consuming and resource-expensive, so CPUs are built with a cache for recently-used addresses: the Translation Lookaside Buffer (TLB). However, the default TLB can only cache a certain number of address mappings. If a requested address mapping is not in the TLB (that is, the TLB is *missed*), the system still needs to read the page table to determine the physical to virtual address mapping.

Because of the relationship between application memory requirements and the size of pages used to cache address mappings, applications with large memory requirements are more likely to suffer performance degradation from TLB misses than applications with minimal memory requirements. It is therefore important to avoid TLB misses wherever possible.

Both HugeTLB and Transparent Huge Page features allow applications to use pages larger than 4 KB. This allows addresses stored in the TLB to reference more memory, which reduces TLB misses and improves application performance.

6.2. MONITORING AND DIAGNOSING PERFORMANCE PROBLEMS

Red Hat Enterprise Linux 7 provides a number of tools that are useful for monitoring system performance

and diagnosing performance problems related to system memory. This section outlines the available tools and gives examples of how to use them to monitor and diagnose memory related performance issues.

6.2.1. Monitoring Memory Usage with vmstat

Vmstat, provided by the procps-ng package, outputs reports on your system's processes, memory, paging, block input/output, interrupts, and CPU activity. It provides an instantaneous report of the average of these events since the machine was last booted, or since the previous report.

The following command displays a table of various event counters and memory statistics.

```
$ vmstat -s
```

For further details of how to use **vmstat**, see Section A.8, “vmstat” or the man page:

```
$ man vmstat
```

6.2.2. Profiling Application Memory Usage with Valgrind

Valgrind is a framework that provides instrumentation to user-space binaries. It ships with a number of tools that can be used to profile and analyze program performance. The **valgrind** tools outlined in this section can help you to detect memory errors such as uninitialized memory use and improper memory allocation or deallocation.

To use **valgrind** or any of its tools, install the valgrind package:

```
# yum install valgrind
```

6.2.2.1. Profiling Memory Usage with Memcheck

Memcheck is the default **valgrind** tool. It detects and reports on a number of memory errors that can be difficult to detect and diagnose, such as:

- Memory access that should not occur
- Undefined or uninitialized value use
- Incorrectly freed heap memory
- Pointer overlap
- Memory leaks

NOTE

Memcheck can only report these errors; it cannot prevent them from occurring. If your program accesses memory in a way that would normally cause a segmentation fault, the segmentation fault still occurs. However, **memcheck** will log an error message immediately prior to the fault.

Because **memcheck** uses instrumentation, applications executed with **memcheck** run ten to thirty times slower than usual.

To run **memcheck** on an application, execute the following command:

```
# valgrind --tool=memcheck application
```

You can also use the following options to focus **memcheck** output on specific types of problem.

--leak-check

After the application finishes executing, **memcheck** searches for memory leaks. The default value is **`--leak-check=summary`**, which prints the number of memory leaks found. You can specify **`--leak-check=yes`** or **`--leak-check=full`** to output details of each individual leak. To disable, specify **`--leak-check=no`**.

--undef-value-errors

The default value is **`--undef-value-errors=yes`**, which reports errors when undefined values are used. You can also specify **`--undef-value-errors=no`**, which will disable this report and slightly speed up Memcheck.

--ignore-ranges

Specifies one or more ranges that **memcheck** should ignore when checking for memory addressability, for example, **`--ignore-ranges=0xPP-0xQQ,0xRR-0xSS`**.

For a full list of **memcheck** options, see the documentation included at **`/usr/share/doc/valgrind-version/valgrind_manual.pdf`**.

6.2.2.2. Profiling Cache Usage with Cachegrind

Cachegrind simulates application interaction with a system's cache hierarchy and branch predictor. It tracks usage of the simulated first level instruction and data caches to detect poor code interaction with this level of cache. It also tracks the last level of cache (second or third level) in order to track memory access. As such, applications executed with **cachegrind** run twenty to one hundred times slower than usual.

Cachegrind gathers statistics for the duration of application execution and outputs a summary to the console. To run **cachegrind** on an application, execute the following command:

```
# valgrind --tool=cachegrind application
```

You can also use the following options to focus **cachegrind** output on a specific problem.

--I1

Specifies the size, associativity, and line size of the first level instruction cache, like so: **`--I1=size,associativity,line_size`**.

--D1

Specifies the size, associativity, and line size of the first level data cache, like so: **`--D1=size,associativity,line_size`**.

--LL

Specifies the size, associativity, and line size of the last level cache, like so: **`--LL=size,associativity,line_size`**.

--cache-sim

Enables or disables the collection of cache access and miss counts. This is enabled (**--cache-sim=yes**) by default. Disabling both this and **--branch-sim** leaves **cachegrind** with no information to collect.

--branch-sim

Enables or disables the collection of branch instruction and incorrect prediction counts. This is enabled (**--branch-sim=yes**) by default. Disabling both this and **--cache-sim** leaves **cachegrind** with no information to collect.

Cachegrind writes detailed profiling information to a per-process **cachegrind.out.*pid*** file, where *pid* is the process identifier. This detailed information can be further processed by the companion **cg_annotate** tool, like so:

```
# cg_annotate cachegrind.out.pid
```

Cachegrind also provides the cg_diff tool, which makes it easier to chart program performance before and after a code change. To compare output files, execute the following command, replacing first with the initial profile output file, and second with the subsequent profile output file.

```
# cg_diff first second
```

The resulting output file can be viewed in greater detail with the **cg_annotate** tool.

For a full list of **cachegrind** options, see the documentation included at **/usr/share/doc/valgrind-version/valgrind_manual.pdf**.

6.2.2.3. Profiling Heap and Stack Space with Massif

Massif measures the heap space used by a specified application. It measures both useful space and any additional space allocated for bookkeeping and alignment purposes. **massif** helps you understand how you can reduce your application's memory use to increase execution speed and reduce the likelihood that your application will exhaust system swap space. Applications executed with **massif** run about twenty times slower than usual.

To run **massif** on an application, execute the following command:

```
# valgrind --tool=massif application
```

You can also use the following options to focus **massif** output on a specific problem.

--heap

Specifies whether **massif** profiles the heap. The default value is **--heap=yes**. Heap profiling can be disabled by setting this to **--heap=no**.

--heap-admin

Specifies the number of bytes per block to use for administration when heap profiling is enabled. The default value is **8** bytes.

--stacks

Specifies whether **massif** profiles the stack. The default value is **--stack=no**, as stack profiling can

greatly slow **massif**. Set this option to **`--stack=yes`** to enable stack profiling. Note that **massif** assumes that the main stack starts with a size of zero in order to better indicate the changes in stack size that relate to the application being profiled.

--time-unit

Specifies the interval at which **massif** gathers profiling data. The default value is **`i`** (instructions executed). You can also specify **`ms`** (milliseconds, or realtime) and **`B`** (bytes allocated or deallocated on the heap and stack). Examining bytes allocated is useful for short run applications and for testing purposes, as it is most reproducible across different hardware.

Massif outputs profiling data to a **`massif.out.`*`pid`*** file, where *pid* is the process identifier of the specified application. The **ms_print** tool graphs this profiling data to show memory consumption over the execution of the application, as well as detailed information about the sites responsible for allocation at points of peak memory allocation. To graph the data from the **`massif.out.`*`pid`*** file, execute the following command:

```
# ms_print massif.out.pid
```

For a full list of **Massif** options, see the documentation included at **`/usr/share/doc/valgrind-version/valgrind_manual.pdf`**.

6.3. CONFIGURING HUGETLB HUGE PAGES

Starting with Red Hat Enterprise Linux 7.1, there are two ways of reserving huge pages: at **`boot time`** and at **`run time`**. Reserving at boot time increases the possibility of success because the memory has not yet been significantly fragmented. However, on NUMA machines, the number of pages is automatically split among NUMA nodes. The run-time method allows you to reserve huge pages per NUMA node. If the run-time reservation is done as early as possible in the boot process, the probability of memory fragmentation is lower.

6.3.1. Configuring Huge Pages at Boot Time

To configure huge pages at boot time, add the following parameters to the kernel boot command line:

hugepages

Defines the number of persistent huge pages configured in the kernel at boot time. The default value is 0. It is only possible to allocate huge pages if there are sufficient physically contiguous free pages in the system. Pages reserved by this parameter cannot be used for other purposes.

This value can be adjusted after boot by changing the value of the **`/proc/sys/vm/nr_hugepages`** file.

In a NUMA system, huge pages assigned with this parameter are divided equally between nodes. You can assign huge pages to specific nodes at runtime by changing the value of the node's **`/sys/devices/system/node/`*`node_id`*`/hugepages/hugepages-1048576kB/nr_hugepages`** file.

For more information, read the relevant kernel documentation, which is installed in **`/usr/share/doc/kernel-doc-kernel_version/Documentation/vm/hugetlbpage.txt`** by default.

hugepagesz

Defines the size of persistent huge pages configured in the kernel at boot time. Valid values are 2 MB and 1 GB. The default value is 2 MB.

default_hugepagesz

Defines the default size of persistent huge pages configured in the kernel at boot time. Valid values are 2 MB and 1 GB. The default value is 2 MB.

Procedure 6.1. Reserving 1 GB Pages During Early Boot

The page size the HugeTLB subsystem supports depends on the architecture. On the AMD64 and Intel 64 architecture, 2 MB huge pages and 1 GB gigantic pages are supported.

1. Create a HugeTLB pool for 1 GB pages by appending the following line to the kernel command-line options:

```
default_hugepagesz=1G hugepagesz=1G
```

2. Create a file named **/usr/lib/systemd/system/hugetlb-gigantic-pages.service** with the following content:

```
[Unit]
Description=HugeTLB Gigantic Pages Reservation
DefaultDependencies=no
Before=dev-hugepages.mount
ConditionPathExists=/sys/devices/system/node
ConditionKernelCommandLine=hugepagesz=1G

[Service]
Type=oneshot
RemainAfterExit=yes
ExecStart=/usr/lib/systemd/hugetlb-reserve-pages

[Install]
WantedBy=sysinit.target
```

3. Create a file named **/usr/lib/systemd/hugetlb-reserve-pages** with the following content:

```
nodes_path=/sys/devices/system/node/
if [ ! -d $nodes_path ]; then
 echo "ERROR: $nodes_path does not exist"
 exit 1
fi

reserve_pages()
{
 echo $1 > $nodes_path/$2/hugepages/hugepages-1048576kB/nr_hugepages
}

# This example reserves 2 1G pages on node0 and 1 1G page on node1.
You
# can modify it to your needs or add more lines to reserve memory in
# other nodes. Don't forget to uncomment the lines, otherwise then
```

```
won't
# be executed.
# reserve_pages 2 node0
# reserve_pages 1 node1
```

4. Modify **/usr/lib/systemd/hugetlb-reserve-pages** according to the comments in the file.

5. Run the following commands to enable early boot reservation:

```
# chmod +x /usr/lib/systemd/hugetlb-reserve-pages
# systemctl enable hugetlb-gigantic-pages
```

NOTE

You can try reserving more 1G pages at runtime by writing to **nr_hugepages** at any time. However, such reservations can fail due to memory fragmention. The most reliable way to reserve 1G pages is by doing it at early boot as described in step 4.

6.3.2. Configuring Huge Pages at Run Time

Use the following parameters to influence huge page behavior at run time:

/sys/devices/system/node/*node_id*/hugepages/hugepages-*size*/nr_hugepages

Defines the number of huge pages of the specified size assigned to the specified NUMA node. This is supported as of Red Hat Enterprise Linux 7.1. The following example moves adds twenty 2048 kB huge pages to **node2**.

```
# numastat -cm | egrep 'Node|Huge'
                 Node 0 Node 1 Node 2 Node 3  Total add
AnonHugePages         0      2      0      8     10
HugePages_Total       0      0      0      0      0
HugePages_Free        0      0      0      0      0
HugePages_Surp        0      0      0      0      0
# echo 20 > /sys/devices/system/node/node2/hugepages/hugepages-
2048kB/nr_hugepages
# numastat -cm | egrep 'Node|Huge'
                 Node 0 Node 1 Node 2 Node 3  Total
AnonHugePages         0      2      0      8     10
HugePages_Total       0      0     40      0     40
HugePages_Free        0      0     40      0     40
HugePages_Surp        0      0      0      0      0
```

/proc/sys/vm/nr_overcommit_hugepages

Defines the maximum number of additional huge pages that can be created and used by the system through overcommitting memory. Writing any non-zero value into this file indicates that the system obtains that number of huge pages from the kernel's normal page pool if the persistent huge page pool is exhausted. As these surplus huge pages become unused, they are then freed and returned to the kernel's normal page pool.

6.4. CONFIGURING TRANSPARENT HUGE PAGES

Transparent Huge Pages (THP) is an alternative solution to HugeTLB. With THP, the kernel automatically assigns huge pages to processes, so huge pages do not need to be reserved manually.

The THP feature has two modes of operation: system-wide and per-process. When THP is enabled system-wide, the kernel tries to assign huge pages to any process when it is possible to allocate huge pages and the process is using a large contiguous virtual memory area. If THP is enabled per-process, the kernel only assigns huge pages to individual processes' memory areas specified with the **`madvise()`** system call.

Note that the THP feature only supports 2-MB pages. Transparent huge pages are enabled by default. To check the current status, run:

```
# cat /sys/kernel/mm/transparent_hugepage/enabled
```

To enable transparent huge pages, run:

```
# echo always > /sys/kernel/mm/transparent_hugepage/enabled
```

To prevent applications from allocating more memory resources than necessary, you can disable huge pages system-wide and only enable them inside MADV_HUGEPAGE madvise regions by running:

```
# echo madvise > /sys/kernel/mm/transparent_hugepage/enabled
```

To disable transparent huge pages, run:

```
# echo never > /sys/kernel/mm/transparent_hugepage/enabled
```

Sometimes, providing low latency to short-lived allocations has higher priority than immediately achieving the best performance with long-lived allocations. In such cases, **`direct compaction`** can be *disabled* while leaving THP enabled.

Direct compaction is a synchronous memory compaction during the huge page allocation. Disabling direct compaction provides no guarantee of saving memory, but can decrease the risk of higher latencies during frequent page faults. Note that if the workload benefits significantly from THP, the performance decreases. To disable direct compaction, run:

```
# echo madvise > /sys/kernel/mm/transparent_hugepage/defrag
```

For comprehensive information on transparent huge pages, see the **`/usr/share/doc/kernel-doc-`*`kernel_version`*`/Documentation/vm/transhuge.txt`** file, which is available after installing the kernel-doc package.

6.5. CONFIGURING SYSTEM MEMORY CAPACITY

This section discusses memory-related kernel parameters that may be useful in improving memory utilization on your system. These parameters can be temporarily set for testing purposes by altering the value of the corresponding file in the **`/proc`** file system. Once you have determined the values that produce optimal performance for your use case, you can set them permanently by using the **`sysctl`** command.

Memory usage is typically configured by setting the value of one or more kernel parameters. These parameters can be set temporarily by altering the contents of files in the /proc file system, or they can be set persistently with the sysctl tool, which is provided by the procps-ng package.

For example, to set the overcommit_memory parameter to 1 temporarily, run the following command:

```
# echo 1 > /proc/sys/vm/overcommit_memory
```

To set this value persistently, add **sysctl vm.overcommit_memory=1** in **/etc/sysctl.conf** then run the following command:

```
# sysctl -p
```

Setting a parameter temporarily is useful for determining the effect the parameter has on your system. You can then set the parameter persistently when you are sure that the parameter's value has the desired effect.

NOTE

To expand your expertise, you might also be interested in the Red Hat Enterprise Linux Performance Tuning (RH442) training course.

6.5.1. Virtual Memory Parameters

The parameters listed in this section are located in **/proc/sys/vm** unless otherwise indicated.

dirty_ratio

A percentage value. When this percentage of total system memory is modified, the system begins writing the modifications to disk with the **pdflush** operation. The default value is **20** percent.

dirty_background_ratio

A percentage value. When this percentage of total system memory is modified, the system begins writing the modifications to disk in the background. The default value is **10** percent.

overcommit_memory

Defines the conditions that determine whether a large memory request is accepted or denied.

The default value is **0**. By default, the kernel performs heuristic memory overcommit handling by estimating the amount of memory available and failing requests that are too large. However, since memory is allocated using a heuristic rather than a precise algorithm, overloading memory is possible with this setting.

When this parameter is set to **1**, the kernel performs no memory overcommit handling. This increases the possibility of memory overload, but improves performance for memory-intensive tasks.

When this parameter is set to **2**, the kernel denies requests for memory equal to or larger than the sum of total available swap space and the percentage of physical RAM specified in ***overcommit_ratio***. This reduces the risk of overcommitting memory, but is recommended only for systems with swap areas larger than their physical memory.

overcommit_ratio

Specifies the percentage of physical RAM considered when ***overcommit_memory*** is set to **2**. The default value is **50**.

max_map_count

Defines the maximum number of memory map areas that a process can use. The default value (**65530**) is appropriate for most cases. Increase this value if your application needs to map more than this number of files.

min_free_kbytes

Specifies the minimum number of kilobytes to keep free across the system. This is used to determine an appropriate value for each low memory zone, each of which is assigned a number of reserved free pages in proportion to their size.

WARNING

Extreme values can damage your system. Setting ***min_free_kbytes*** to an extremely low value prevents the system from reclaiming memory, which can result in system hangs and OOM-killing processes. However, setting ***min_free_kbytes*** too high (for example, to 5–10% of total system memory) causes the system to enter an out-of-memory state immediately, resulting in the system spending too much time reclaiming memory.

oom_adj

In the event that the system runs out of memory and the ***panic_on_oom*** parameter is set to **0**, the **oom_killer** function kills processes until the system can recover, starting from the process with the highest **oom_score**.

The ***oom_adj*** parameter helps determine the ***oom_score*** of a process. This parameter is set per process identifier. A value of **-17** disables the **oom_killer** for that process. Other valid values are from **-16** to **15**.

NOTE

Processes spawned by an adjusted process inherit the **oom_score** of the process.

swappiness

The swappiness value, ranging from **0** to **100**, controls the degree to which the system favors anonymous memory or the page cache. A high value improves file-system performance while aggressively swapping less active processes out of RAM. A low value avoids swapping processes out of memory, which usually decreases latency at the cost of I/O performance. The default value is **60**.

WARNING

Setting **swappiness==0** will very aggressively avoids swapping out, which increase the risk of OOM killing under strong memory and I/O pressure.

6.5.2. File System Parameters

Parameters listed in this section are located in **/proc/sys/fs** unless otherwise indicated.

aio-max-nr

Defines the maximum allowed number of events in all active asynchronous input/output contexts. The default value is **65536**. Modifying this value does not pre-allocate or resize any kernel data structures.

file-max

Determines the maximum number of file handles for the entire system. The default value on Red Hat Enterprise Linux 7 is the maximum of either **8192**, or one tenth of the free memory pages available at the time the kernel starts.

Raising this value can resolve errors caused by a lack of available file handles.

6.5.3. Kernel Parameters

Default values for the following parameters, located in the **/proc/sys/kernel/** directory, can be calculated by the kernel at boot time depending on available system resources.

msgmax

Defines the maximum allowable size in bytes of any single message in a message queue. This value must not exceed the size of the queue (***msgmnb***). To determine the current ***msgmax*** value on your system, use:

```
# sysctl kernel.msgmax
```

msgmnb

Defines the maximum size in bytes of a single message queue. To determine the current ***msgmnb*** value on your system, use:

```
# sysctl kernel.msgmnb
```

msgmni

Defines the maximum number of message queue identifiers, and therefore the maximum number of queues. To determine the current ***msgmni*** value on your system, use:

```
# sysctl kernel.msgmni
```

shmall

Defines the total amount of shared memory (in bytes) that can be used on the system at one time. To determine the current ***shmall*** value on your system, use:

```
# sysctl kernel.shmall
```

shmmax

Defines the maximum size (in bytes) of a single shared memory segment allowed by the kernel. To determine the current ***shmmax*** value on your system, use:

```
# sysctl kernel.shmmax
```

shmmni

Defines the system-wide maximum number of shared memory segments. The default value is **4096** on all systems.

threads-max

Defines the system-wide maximum number of threads available to the kernel at one time. To determine the current ***threads-max*** value on your system, use:

```
# sysctl kernel.threads-max
```

The default value is the result of:

```
mempages / (8 * THREAD_SIZE / PAGE SIZE )
```

The minimum value is **20**.

CHAPTER 7. STORAGE AND FILE SYSTEMS

This chapter outlines supported file systems and configuration options that affect application performance for both I/O and file systems in Red Hat Enterprise Linux 7. Section 7.1, “Considerations” discusses the I/O and file system related factors that affect performance. Section 7.2, “Monitoring and Diagnosing Performance Problems” teaches you how to use Red Hat Enterprise Linux 7 tools to diagnose performance problems related to I/O or file system configuration details. Section 7.4, “Configuration Tools” discusses the tools and strategies you can use to solve I/O and file system related performance problems in Red Hat Enterprise Linux 7.

7.1. CONSIDERATIONS

The appropriate settings for storage and file system performance are highly dependent on the purpose of the storage. I/O and file system performance can be affected by any of the following factors:

- Data write or read patterns
- Data alignment with underlying geometry
- Block size
- File system size
- Journal size and location
- Recording access times
- Ensuring data reliability
- Pre-fetching data
- Pre-allocating disk space
- File fragmentation
- Resource contention

Read this chapter to gain an understanding of the formatting and mount options that affect file system throughput, scalability, responsiveness, resource usage, and availability.

7.1.1. I/O Schedulers

The I/O scheduler determines when and for how long I/O operations run on a storage device. It is also known as the I/O elevator.

Red Hat Enterprise Linux 7 provides three I/O schedulers.

deadline

The default I/O scheduler for all block devices, except for SATA disks. **`Deadline`** attempts to provide a guaranteed latency for requests from the point at which requests reach the I/O scheduler. This scheduler is suitable for most use cases, but particularly those in which read operations occur more often than write operations.

Queued I/O requests are sorted into a read or write batch and then scheduled for execution in increasing LBA order. Read batches take precedence over write batches by default, as applications are more likely to block on read I/O. After a batch is processed, **`deadline`** checks how long write

operations have been starved of processor time and schedules the next read or write batch as appropriate. The number of requests to handle per batch, the number of read batches to issue per write batch, and the amount of time before requests expire are all configurable; see Section 7.4.4, “Tuning the Deadline Scheduler” for details.

cfq

The default scheduler only for devices identified as SATA disks. The Completely Fair Queueing scheduler, **cfq**, divides processes into three separate classes: real time, best effort, and idle. Processes in the real time class are always performed before processes in the best effort class, which are always performed before processes in the idle class. This means that processes in the real time class can starve both best effort and idle processes of processor time. Processes are assigned to the best effort class by default.

cfq uses historical data to anticipate whether an application will issue more I/O requests in the near future. If more I/O is expected, **cfq** idles to wait for the new I/O, even if I/O from other processes is waiting to be processed.

Because of this tendency to idle, the cfq scheduler should not be used in conjunction with hardware that does not incur a large seek penalty unless it is tuned for this purpose. It should also not be used in conjunction with other non-work-conserving schedulers, such as a host-based hardware RAID controller, as stacking these schedulers tends to cause a large amount of latency.

cfq behavior is highly configurable; see Section 7.4.5, “Tuning the CFQ Scheduler” for details.

noop

The **noop** I/O scheduler implements a simple FIFO (first-in first-out) scheduling algorithm. Requests are merged at the generic block layer through a simple last-hit cache. This can be the best scheduler for CPU-bound systems using fast storage.

For details on setting a different default I/O scheduler, or specifying a different scheduler for a particular device, see Section 7.4, “Configuration Tools”.

7.1.2. File Systems

Read this section for details about supported file systems in Red Hat Enterprise Linux 7, their recommended use cases, and the format and mount options available to file systems in general. Detailed tuning recommendations for these file systems are available in Section 7.4.7, “Configuring File Systems for Performance”.

7.1.2.1. XFS

XFS is a robust and highly scalable 64-bit file system. It is the default file system in Red Hat Enterprise Linux 7. XFS uses extent-based allocation, and features a number of allocation schemes, including pre-allocation and delayed allocation, both of which reduce fragmentation and aid performance. It also supports metadata journaling, which can facilitate crash recovery. XFS can be defragmented and enlarged while mounted and active, and Red Hat Enterprise Linux 7 supports several XFS-specific backup and restore utilities.

As of Red Hat Enterprise Linux 7.0 GA, XFS is supported to a maximum file system size of 500 TB, and a maximum file offset of 8 EB (sparse files). For details about administering XFS, see the *Red Hat Enterprise Linux 7 Storage Administration Guide*, available from http://access.redhat.com/site/documentation/Red_Hat_Enterprise_Linux/. For assistance tuning XFS for a specific purpose, see Section 7.4.7.1, “Tuning XFS”.

7.1.2.2. Ext4

Ext4 is a scalable extension of the ext3 file system. Its default behavior is optimal for most work loads. However, it is supported only to a maximum file system size of 50 TB, and a maximum file size of 16 TB. For details about administering ext4, see the *Red Hat Enterprise Linux 7 Storage Administration Guide*, available from http://access.redhat.com/site/documentation/Red_Hat_Enterprise_Linux/. For assistance tuning ext4 for a specific purpose, see Section 7.4.7.2, “Tuning ext4”.

7.1.2.3. Btrfs (Technology Preview)

The default file system for Red Hat Enterprise Linux 7 is XFS. Btrfs (B-tree file system), a relatively new copy-on-write (COW) file system, is shipped as a Technology Preview. Some of the unique Btrfs features include:

- The ability to take snapshots of specific files, volumes or sub-volumes rather than the whole file system;
- supporting several versions of redundant array of inexpensive disks (RAID);
- back referencing map I/O errors to file system objects;
- transparent compression (all files on the partition are automatically compressed);
- checksums on data and meta-data.

Although Btrfs is considered a stable file system, it is under constant development, so some functionality, such as the repair tools, are basic compared to more mature file systems.

Currently, selecting Btrfs is suitable when advanced features (such as snapshots, compression, and file data checksums) are required, but performance is relatively unimportant. If advanced features are not required, the risk of failure and comparably weak performance over time make other file systems preferable. Another drawback, compared to other file systems, is the maximum supported file system size of 50 TB.

For more information, see Section 7.4.7.3, “Tuning Btrfs”, and the chapter on Btrfs in the Red Hat Enterprise Linux 7 Storage Administration Guide.

7.1.2.4. GFS2

Global File System 2 (GFS2) is part of the High Availability Add-On that provides clustered file system support to Red Hat Enterprise Linux 7. GFS2 provides a consistent file system image across all servers in a cluster, which allows servers to read from and write to a single shared file system.

GFS2 is supported to a maximum file system size of 100 TB.

For details on administering GFS2, see the Red Hat Enterprise Linux 7 Storage Administration Guide. For information on tuning GFS2 for a specific purpose, see Section 7.4.7.4, “Tuning GFS2”.

7.1.3. Generic Tuning Considerations for File Systems

This section covers tuning considerations common to all file systems. For tuning recommendations specific to your file system, see Section 7.4.7, “Configuring File Systems for Performance”.

7.1.3.1. Considerations at Format Time

Some file system configuration decisions cannot be changed after the device is formatted. This section covers the options available to you for decisions that must be made before you format your storage device.

Size

Create an appropriately-sized file system for your workload. Smaller file systems have proportionally shorter backup times and require less time and memory for file system checks. However, if your file system is too small, its performance will suffer from high fragmentation.

Block size

The block is the unit of work for the file system. The block size determines how much data can be stored in a single block, and therefore the smallest amount of data that is written or read at one time.

The default block size is appropriate for most use cases. However, your file system will perform better and store data more efficiently if the block size (or the size of multiple blocks) is the same as or slightly larger than amount of data that is typically read or written at one time. A small file will still use an entire block. Files can be spread across multiple blocks, but this can create additional runtime overhead. Additionally, some file systems are limited to a certain number of blocks, which in turn limits the maximum size of the file system.

Block size is specified as part of the file system options when formatting a device with the **mkfs** command. The parameter that specifies the block size varies with the file system; see the **mkfs** man page for your file system for details. For example, to see the options available when formatting an XFS file system, execute the following command.

```
$ man mkfs.xfs
```

Geometry

File system geometry is concerned with the distribution of data across a file system. If your system uses striped storage, like RAID, you can improve performance by aligning data and metadata with the underlying storage geometry when you format the device.

Many devices export recommended geometry, which is then set automatically when the devices are formatted with a particular file system. If your device does not export these recommendations, or you want to change the recommended settings, you must specify geometry manually when you format the device with **mkfs**.

The parameters that specify file system geometry vary with the file system; see the **mkfs** man page for your file system for details. For example, to see the options available when formatting an ext4 file system, execute the following command.

```
$ man mkfs.ext4
```

External journals

Journaling file systems document the changes that will be made during a write operation in a journal file prior to the operation being executed. This reduces the likelihood that a storage device will become corrupted in the event of a system crash or power failure, and speeds up the recovery process.

Metadata-intensive workloads involve very frequent updates to the journal. A larger journal uses more memory, but reduces the frequency of write operations. Additionally, you can improve the seek time of a device with a metadata-intensive workload by placing its journal on dedicated storage that is as fast as, or faster than, the primary storage.

WARNING

Ensure that external journals are reliable. Losing an external journal device will cause file system corruption.

External journals must be created at format time, with journal devices being specified at mount time. For details, see the **mkfs** and **mount** man pages.

```
$ man mkfs
```

```
$ man mount
```

7.1.3.2. Considerations at Mount Time

This section covers tuning decisions that apply to most file systems and can be specified as the device is mounted.

Barriers

File system barriers ensure that file system metadata is correctly written and ordered on persistent storage, and that data transmitted with **fsync** persists across a power outage. On previous versions of Red Hat Enterprise Linux, enabling file system barriers could significantly slow applications that relied heavily on **fsync**, or created and deleted many small files.

In Red Hat Enterprise Linux 7, file system barrier performance has been improved such that the performance effects of disabling file system barriers are negligible (less than 3%).

For further information, see the *Red Hat Enterprise Linux 7 Storage Administration Guide*, available from http://access.redhat.com/site/documentation/Red_Hat_Enterprise_Linux/.

Access Time

Every time a file is read, its metadata is updated with the time at which access occurred (**atime**). This involves additional write I/O. In most cases, this overhead is minimal, as by default Red Hat Enterprise Linux 7 updates the **atime** field only when the previous access time was older than the times of last modification (**mtime**) or status change (**ctime**).

However, if updating this metadata is time consuming, and if accurate access time data is not required, you can mount the file system with the **noatime** mount option. This disables updates to metadata when a file is read. It also enables **nodiratime** behavior, which disables updates to metadata when a directory is read.

Read-ahead

Read-ahead behavior speeds up file access by pre-fetching data that is likely to be needed soon and loading it into the page cache, where it can be retrieved more quickly than if it were on disk. The higher the read-ahead value, the further ahead the system pre-fetches data.

Red Hat Enterprise Linux attempts to set an appropriate read-ahead value based on what it detects about your file system. However, accurate detection is not always possible. For example, if a storage

array presents itself to the system as a single LUN, the system detects the single LUN, and does not set the appropriate read-ahead value for an array.

Workloads that involve heavy streaming of sequential I/O often benefit from high read-ahead values. The storage-related tuned profiles provided with Red Hat Enterprise Linux 7 raise the read-ahead value, as does using LVM striping, but these adjustments are not always sufficient for all workloads.

The parameters that define read-ahead behavior vary with the file system; see the mount man page for details.

```
$ man mount
```

7.1.3.3. Maintenance

Regularly discarding blocks that are not in use by the file system is a recommended practice for both solid-state disks and thinly-provisioned storage. There are two methods of discarding unused blocks: batch discard and online discard.

Batch discard

This type of discard is part of the **fstrim** command. It discards all unused blocks in a file system that match criteria specified by the administrator.

Red Hat Enterprise Linux 7 supports batch discard on XFS and ext4 formatted devices that support physical discard operations (that is, on HDD devices where the value of **/sys/block/*devname*/queue/discard_max_bytes** is not zero, and SSD devices where the value of **/sys/block/*devname*/queue/discard_granularity** is not **0**).

Online discard

This type of discard operation is configured at mount time with the **discard** option, and runs in real time without user intervention. However, online discard only discards blocks that are transitioning from used to free. Red Hat Enterprise Linux 7 supports online discard on XFS and ext4 formatted devices.

Red Hat recommends batch discard except where online discard is required to maintain performance, or where batch discard is not feasible for the system's workload.

Pre-allocation

Pre-allocation marks disk space as being allocated to a file without writing any data into that space. This can be useful in limiting data fragmentation and poor read performance. Red Hat Enterprise Linux 7 supports pre-allocating space on XFS, ext4, and GFS2 devices at mount time; see the **mount** man page for the appropriate parameter for your file system. Applications can also benefit from pre-allocating space by using the **fallocate(2) glibc** call.

7.2. MONITORING AND DIAGNOSING PERFORMANCE PROBLEMS

Red Hat Enterprise Linux 7 provides a number of tools that are useful for monitoring system performance and diagnosing performance problems related to I/O and file systems and their configuration. This section outlines the available tools and gives examples of how to use them to monitor and diagnose I/O and file system related performance issues.

7.2.1. Monitoring System Performance with vmstat

Vmstat reports on processes, memory, paging, block I/O, interrupts, and CPU activity across the entire system. It can help administrators determine whether the I/O subsystem is responsible for any performance issues.

The information most relevant to I/O performance is in the following columns:

si

Swap in, or reads from swap space, in KB.

so

Swap out, or writes to swap space, in KB.

bi

Block in, or block write operations, in KB.

bo

Block out, or block read operations, in KB.

wa

The portion of the queue that is waiting for I/O operations to complete.

Swap in and swap out are particularly useful when your swap space and your data are on the same device, and as indicators of memory usage.

Additionally, the free, buff, and cache columns can help identify write-back frequency. A sudden drop in cache values and an increase in free values indicates that write-back and page cache invalidation has begun.

If analysis with **vmstat** shows that the I/O subsystem is responsible for reduced performance, administrators can use **iostat** to determine the responsible I/O device.

vmstat is provided by the procps-ng package. For detailed information about using **vmstat**, see the man page:

```
$ man vmstat
```

7.2.2. Monitoring I/O Performance with iostat

Iostat is provided by the sysstat package. It reports on I/O device load in your system. If analysis with **vmstat** shows that the I/O subsystem is responsible for reduced performance, you can use **iostat** to determine the I/O device responsible.

You can focus the output of **iostat** reports on a specific device by using the parameters defined in the **iostat** man page:

```
$ man iostat
```

7.2.2.1. Detailed I/O Analysis with blktrace

Blktrace provides detailed information about how time is spent in the I/O subsystem. The companion utility **blkparse** reads the raw output from **blktrace** and produces a human readable summary of input and output operations recorded by **blktrace**.

For more detailed information about this tool, see the man page:

```
$ man blktrace
```

```
$ man blkparse
```

7.2.2.2. Analyzing blktrace Output with btt

Btt is provided as part of the blktrace package. It analyzes **blktrace** output and displays the amount of time that data spends in each area of the I/O stack, making it easier to spot bottlenecks in the I/O subsystem.

For example, if **btt** shows that the time between requests being sent to the block layer (**Q2Q**) is larger than the total time that requests spent in the block layer (**Q2C**), the I/O subsystem may not be responsible for performance issues. If the device takes a long time to service a request (**D2C**), the device may be overloaded, or the workload sent to the device may be sub-optimal. If block I/O is queued for a long time before being assigned a request (**Q2G**), it may indicate that the storage in use is unable to serve the I/O load.

For more detailed information about this tool, see the man page:

```
$ man btt
```

7.2.2.3. Analyzing blktrace Output with seekwatcher

The **seekwatcher** tool can use **blktrace** output to graph I/O over time. It focuses on the Logical Block Address (LBA) of disk I/O, throughput in megabytes per second, the number of seeks per second, and I/O operations per second. This can help to identify when you are hitting the operations-per-second limit of a device.

For more detailed information about this tool, see the man page:

```
$ man seekwatcher
```

7.2.3. Storage Monitoring with SystemTap

The *Red Hat Enterprise Linux 7 SystemTap Beginner's Guide* includes several sample scripts that are useful for profiling and monitoring storage performance.

The following **SystemTap** example scripts relate to storage performance and may be useful in diagnosing storage or file system performance problems. By default they are installed to the **/usr/share/doc/systemtap-client/examples/io** directory.

disktop.stp

Checks the status of reading/writing disk every 5 seconds and outputs the top ten entries during that period.

iotime.stp

Prints the amount of time spent on read and write operations, and the number of bytes read and written.

`traceio.stp`

Prints the top ten executables based on cumulative I/O traffic observed, every second.

`traceio2.stp`

Prints the executable name and process identifier as reads and writes to the specified device occur.

`inodewatch.stp`

Prints the executable name and process identifier each time a read or write occurs to the specified inode on the specified major/minor device.

`inodewatch2.stp`

Prints the executable name, process identifier, and attributes each time the attributes are changed on the specified inode on the specified major/minor device.

The *Red Hat Enterprise Linux 7 SystemTap Beginner's Guide* is available from http://access.redhat.com/site/documentation/Red_Hat_Enterprise_Linux/.

7.3. SOLID-STATE DISKS

Solid-state disks (SSD) use NAND flash chips rather than rotating magnetic platters to store persistent data. They provide a constant access time for data across their full Logical Block Address range, and do not incur measurable seek costs like their rotating counterparts. They are more expensive per gigabyte of storage space and have a lesser storage density, but they also have lower latency and greater throughput than HDDs.

Performance generally degrades as the used blocks on an SSD approach the capacity of the disk. The degree of degradation varies by vendor, but all devices experience degradation in this circumstance. Enabling discard behavior can help to alleviate this degradation. For details, see Section 7.1.3.3, “Maintenance”.

The default I/O scheduler and virtual memory options are suitable for use with SSDs.

For more information on SSD, see the Solid-State Disk Deployment Guidelines chapter in the *Red Hat Enterprise Linux 7 Storage Administration Guide*.

SSD Tuning Considerations

Consider the following factors when configuring settings that can affect SSD performance:

I/O Scheduler

Any I/O scheduler is expected to perform well with most SSDs. However, as with any other storage type, Red Hat recommends benchmarking to determine the optimal configuration for a given workload. When using SSDs, Red Hat advises changing the I/O scheduler only for benchmarking particular workloads. For instructions on how to switch between I/O schedulers, see the **`/usr/share/doc/kernel-version/Documentation/block/switching-sched.txt`** file.

As of Red Hat Enterprise Linuxnbsp 7.0, the default I/O scheduler is Deadline, except for use with SATA drives, which use CFQ as the default I/O scheduler. For faster storage, Deadline can outperform CFQ leading to better I/O performance without the need for specific tuning. Sometimes, the default is not

suitable for certain disks, such as SAS rotational disks. In such cases, change the I/O scheduler to CFQ.

Virtual Memory

Like the I/O scheduler, virtual memory (VM) subsystem requires no special tuning. Given the fast nature of I/O on SSD, try turning down the **vm_dirty_background_ratio** and **vm_dirty_ratio** settings, as increased write-out activity does not usually have a negative impact on the latency of other operations on the disk. However, this tuning can generate *more overall I/O*, and is therefore not generally recommended without workload-specific testing.

Swap

An SSD can also be used as a swap device, and is likely to produce good page-out and page-in performance.

7.4. CONFIGURATION TOOLS

Red Hat Enterprise Linux provides a number of tools to assist administrators in configuring the storage and file systems. This section outlines the available tools and provides examples of how they can be used to solve I/O and file system related performance problems in Red Hat Enterprise Linux 7.

7.4.1. Configuring Tuning Profiles for Storage Performance

The **Tuned** service provides a number of profiles designed to improve performance for specific use cases. The following profiles are particularly useful for improving storage performance.

- latency-performance
- throughput-performance (the default)

To configure a profile on your system, run the following command, replacing *name* with the name of the profile you want to use.

```
$ tuned-adm profile name
```

The **tuned-adm recommend** command recommends an appropriate profile for your system.

For further details about these profiles or additional configuration options, see Section A.5, “tuned-adm”.

7.4.2. Setting the Default I/O Scheduler

The default I/O scheduler is the scheduler that is used if no other scheduler is explicitly specified for the device.

If no default scheduler is specified, the **cfq** scheduler is used for SATA drives, and the **deadline** scheduler is used for all other drives. If you specify a default scheduler by following the instructions in this section, that default scheduler is applied to all devices.

To set the default I/O scheduler, you can use the **Tuned** tool, or modify the **/etc/grub/default** file manually.

Red Hat recommends using the **Tuned** tool to specify the default I/O scheduler on a booted system. To set the ***elevator*** parameter, enable the **disk** plug-in. For information on the **disk** plug-in, see Section 3.1.1, “Plug-ins” in the *Tuned* chapter.

To modify the default scheduler by using GRUB 2, append the ***elevator*** parameter to the kernel command line, either at boot time, or when the system is booted. You can use the **Tuned** tool, or modify the **/etc/grub/default** file manually, as described in Procedure 7.1, “Setting the Default I/O Scheduler by Using GRUB 2”.

Procedure 7.1. Setting the Default I/O Scheduler by Using GRUB 2

To set the default I/O Scheduler on a booted system and make the configuration persist after reboot:

1. Add the ***elevator*** parameter to the **GRUB_CMDLINE_LINUX** line in the **/etc/grub/default** file.

```
# cat /etc/default/grub
...
GRUB_CMDLINE_LINUX="crashkernel=auto rd.lvm.lv=vg00/lvroot
rd.lvm.lv=vg00/lvswap elevator=noop"
...
```

   In Red Hat Enterprise Linux 7, the available schedulers are **deadline**, **noop**, and **cfq**. For more information, see the **cfq-iosched.txt** and **deadline-iosched.txt** files in the documentation for your kernel, available after installing the kernel-doc package.

2. Create a new configuration with the ***elevator*** parameter added.

   The location of the GRUB 2 configuration file is different on systems with the BIOS firmware and on systems with UEFI. Use one of the following commands to recreate the GRUB 2 configuration file.

   - On a system with the BIOS firmware, use:

```
# grub2-mkconfig -o /boot/grub2/grub.cfg
```

   - On a system with the UEFI firmware, use:

```
# grub2-mkconfig -o /boot/efi/EFI/redhat/grub.cfg
```

3. Reboot the system for the change to take effect.

   For more information on version 2 of the GNU GRand Unified Bootloader (GRUB 2), see the Working with the GRUB 2 Boot Loader chapter of the *Red Hat Enterprise Linux 7 System Administrator's Guide*.

7.4.3. Generic Block Device Tuning Parameters

The generic tuning parameters listed in this section are available within the **/sys/block/sd*X*/queue/** directory. The listed tuning parameters are separate from I/O scheduler tuning, and are applicable to all I/O schedulers.

add_random

Some I/O events contribute to the entropy pool for **/dev/random**. This parameter can be set to **0** if the overhead of these contributions becomes measurable.

iostats

The default value is **1** (***enabled***). Setting ***iostats*** to **0** disables the gathering of I/O statistics for the device, which removes a small amount of overhead with the I/O path. Setting ***iostats*** to **0** might slightly improve performance for very high performance devices, such as certain NVMe solid-state storage devices. It is recommended to to leave ***iostats*** enabled unless otherwise specified for the given storage model by the vendor.

If you disable ***iostats***, the I/O statistics for the device are no longer present within the **/proc/diskstats** file. The content of **/sys/diskstats** is the source of I/O information for monitoring I/O tools, such as **sar** or ***iostats***. Therefore, if you disable the ***iostats*** parameter for a device, the device is no longer present in the output of I/O monitoring tools.

max_sectors_kb

Specifies the maximum size of an I/O request in kilobytes. The default value is **512** KB. The minimum value for this parameter is determined by the logical block size of the storage device. The maximum value for this parameter is determined by the value of ***max_hw_sectors_kb***.

Certain solid-state disks perform poorly when the I/O requests are larger than the internal erase block size. To determine if this is the case of the solid-state disk model attached to the system, check with the hardware vendor, and follow their recommendations. Red Hat recommends ***max_sectors_kb*** to always be a multiple of the optimal I/O size and the internal erase block size. Use a value of ***logical_block_size*** for either parameter if they are zero or not specified by the storage device.

nomerges

Most workloads benefit from request merging. However, disabling merges can be useful for debugging purposes. By default, the ***nomerges*** parameter is set to **0**, which enables merging. To disable simple one-hit merging, set ***nomerges*** to **1**. To disable all types of merging, set ***nomerges*** to **2**.

nr_requests

Specifies the maximum number of read and write requests that can be queued at one time. The default value is **128**, which means that 128 read requests and 128 write requests can be queued before the next process to request a read or write is put to sleep.

For latency-sensitive applications, lower the value of this parameter and limit the command queue depth on the storage so that write-back I/O cannot fill the device queue with write requests. When the device queue fills, other processes attempting to perform I/O operations are put to sleep until queue space becomes available. Requests are then allocated in a round-robin manner, which prevents one process from continuously consuming all spots in the queue.

The maximum number of I/O operations within the I/O scheduler is **nr_requests*2**. As stated, ***nr_requests*** is applied separately for reads and writes. Note that ***nr_requests*** only applies to the I/O operations within the I/O scheduler and not to I/O operations already dispatched to the underlying device. Therefore, the maximum outstanding limit of I/O operations against a device is **(nr_requests*2)+(queue_depth)** where ***queue_depth*** is **/sys/block/sdN/device/queue_depth**, sometimes also referred to as the LUN queue depth. You can see this total outstanding number of I/O operations in, for example, the output of **iostat** in the **avgqu-sz** column.

optimal_io_size

Some storage devices report an optimal I/O size through this parameter. If this value is reported, Red Hat recommends that applications issue I/O aligned to and in multiples of the optimal I/O size wherever possible.

read_ahead_kb

Defines the maximum number of kilobytes that the operating system may read ahead during a sequential read operation. As a result, the likely-needed information is already present within the kernel page cache for the next sequential read, which improves read I/O performance.

Device mappers often benefit from a high ***read_ahead_kb*** value. 128 KB for each device to be mapped is a good starting point, but increasing the ***read_ahead_kb*** value up to 4–8 MB might improve performance in application environments where sequential reading of large files takes place.

rotational

Some solid-state disks do not correctly advertise their solid-state status, and are mounted as traditional rotational disks. If your solid-state device does does not set this to **0** automatically, set it manually to disable unnecessary seek-reducing logic in the scheduler.

rq_affinity

By default, I/O completions can be processed on a different processor than the processor that issued the I/O request. Set ***rq_affinity*** to **1** to disable this ability and perform completions only on the processor that issued the I/O request. This can improve the effectiveness of processor data caching.

scheduler

To set the scheduler or scheduler preference order for a particular storage device, edit the **/sys/block/*devname*/queue/scheduler** file, where *devname* is the name of the device you want to configure.

```
# echo cfq > /sys/block/hda/queue/scheduler
```

7.4.4. Tuning the Deadline Scheduler

When **deadline** is in use, queued I/O requests are sorted into a read or write batch and then scheduled for execution in increasing LBA order. Read batches take precedence over write batches by default, as applications are more likely to block on read I/O. After a batch is processed, **deadline** checks how long write operations have been starved of processor time and schedules the next read or write batch as appropriate.

The following parameters affect the behavior of the **deadline** scheduler.

fifo_batch

The number of read or write operations to issue in a single batch. The default value is **16**. A higher value can increase throughput, but will also increase latency.

front_merges

If your workload will never generate front merges, this tunable can be set to **0**. However, unless you have measured the overhead of this check, Red Hat recommends the default value of **1**.

read_expire

The number of milliseconds in which a read request should be scheduled for service. The default value is **500** (0.5 seconds).

write_expire

The number of milliseconds in which a write request should be scheduled for service. The default value is **5000** (5 seconds).

writes_starved

The number of read batches that can be processed before processing a write batch. The higher this value is set, the greater the preference given to read batches.

7.4.5. Tuning the CFQ Scheduler

When CFQ is in use, processes are placed into three classes: real time, best effort, and idle. All real time processes are scheduled before any best effort processes, which are scheduled before any idle processes. By default, processes are classed as best effort. You can manually adjust the class of a process with the **ionice** command.

You can further adjust the behavior of the CFQ scheduler with the following parameters. These parameters are set on a per-device basis by altering the specified files under the **/sys/block/*devname*/queue/iosched** directory.

back_seek_max

The maximum distance in kilobytes that CFQ will perform a backward seek. The default value is **16** KB. Backward seeks typically damage performance, so large values are not recommended.

back_seek_penalty

The multiplier applied to backward seeks when the disk head is deciding whether to move forward or backward. The default value is **2**. If the disk head position is at 1024 KB, and there are equidistant requests in the system (1008 KB and 1040 KB, for example), the ***back_seek_penalty*** is applied to backward seek distances and the disk moves forward.

fifo_expire_async

The length of time in milliseconds that an asynchronous (buffered write) request can remain unserviced. After this amount of time expires, a single starved asynchronous request is moved to the dispatch list. The default value is **250** milliseconds.

fifo_expire_sync

The length of time in milliseconds that a synchronous (read or **O_DIRECT** write) request can remain unserviced. After this amount of time expires, a single starved synchronous request is moved to the dispatch list. The default value is **125** milliseconds.

group_idle

This parameter is set to **0** (disabled) by default. When set to **1** (enabled), the **cfq** scheduler idles on the last process that is issuing I/O in a control group. This is useful when using proportional weight I/O control groups and when ***slice_idle*** is set to **0** (on fast storage).

group_isolation

This parameter is set to **0** (disabled) by default. When set to **1** (enabled), it provides stronger isolation between groups, but reduces throughput, as fairness is applied to both random and sequential workloads. When ***group_isolation*** is disabled (set to **0**), fairness is provided to sequential workloads only. For more information, see the installed documentation in **/usr/share/doc/kernel-doc-version/Documentation/cgroups/blkio-controller.txt**.

low_latency

This parameter is set to **1** (enabled) by default. When enabled, **cfq** favors fairness over throughput by providing a maximum wait time of **300** ms for each process issuing I/O on a device. When this parameter is set to **0** (disabled), target latency is ignored and each process receives a full time slice.

quantum

This parameter defines the number of I/O requests that **cfq** sends to one device at one time, essentially limiting queue depth. The default value is **8** requests. The device being used may support greater queue depth, but increasing the value of quantum will also increase latency, especially for large sequential write work loads.

slice_async

This parameter defines the length of the time slice (in milliseconds) allotted to each process issuing asynchronous I/O requests. The default value is **40** milliseconds.

slice_idle

This parameter specifies the length of time in milliseconds that cfq idles while waiting for further requests. The default value is **0** (no idling at the queue or service tree level). The default value is ideal for throughput on external RAID storage, but can degrade throughput on internal non-RAID storage as it increases the overall number of seek operations.

slice_sync

This parameter defines the length of the time slice (in milliseconds) allotted to each process issuing synchronous I/O requests. The default value is **100** ms.

7.4.5.1. Tuning CFQ for Fast Storage

The **cfq** scheduler is not recommended for hardware that does not suffer a large seek penalty, such as fast external storage arrays or solid-state disks. If your use case requires **cfq** to be used on this storage, you will need to edit the following configuration files:

- Set **/sys/block/*devname*/queue/iosched/slice_idle** to **0**
- Set **/sys/block/*devname*/queue/iosched/quantum** to **64**
- Set **/sys/block/*devname*/queue/iosched/group_idle** to **1**

7.4.6. Tuning the noop Scheduler

The **noop** I/O scheduler is primarily useful for CPU-bound systems that use fast storage. Also, the **noop** I/O scheduler is commonly, but not exclusively, used on virtual machines when they are performing I/O operations to virtual disks.

There are no tunable parameters specific to the **noop** I/O scheduler.

7.4.7. Configuring File Systems for Performance

This section covers the tuning parameters specific to each file system supported in Red Hat Enterprise Linux 7. Parameters are divided according to whether their values should be configured when you format the storage device, or when you mount the formatted device.

Where loss in performance is caused by file fragmentation or resource contention, performance can generally be improved by reconfiguring the file system. However, in some cases the application may need to be altered. In this case, Red Hat recommends contacting Customer Support for assistance.

7.4.7.1. Tuning XFS

This section covers some of the tuning parameters available to XFS file systems at format and at mount time.

The default formatting and mount settings for XFS are suitable for most workloads. Red Hat recommends changing them only if specific configuration changes are expected to benefit your workload.

7.4.7.1.1. Formatting Options

For further details about any of these formatting options, see the man page:

```
$ man mkfs.xfs
```

Directory block size

The directory block size affects the amount of directory information that can be retrieved or modified per I/O operation. The minimum value for directory block size is the file system block size (4 KB by default). The maximum value for directory block size is **64** KB.

At a given directory block size, a larger directory requires more I/O than a smaller directory. A system with a larger directory block size also consumes more processing power per I/O operation than a system with a smaller directory block size. It is therefore recommended to have as small a directory and directory block size as possible for your workload.

Red Hat recommends the directory block sizes listed in Table 7.1, “Recommended Maximum Directory Entries for Directory Block Sizes” for file systems with no more than the listed number of entries for write-heavy and read-heavy workloads.

Table 7.1. Recommended Maximum Directory Entries for Directory Block Sizes

Directory block size	Max. entries (read-heavy)	Max. entries (write-heavy)
4 KB	100,000–200,000	1,000,000–2,000,000
16 KB	100,000–1,000,000	1,000,000–10,000,000
64 KB	>1,000,000	>10,000,000

For detailed information about the effect of directory block size on read and write workloads in file systems of different sizes, see the XFS documentation.

To configure directory block size, use the **`mkfs.xfs -l`** option. See the **`mkfs.xfs`** man page for details.

Allocation groups

An allocation group is an independent structure that indexes free space and allocated inodes across a section of the file system. Each allocation group can be modified independently, allowing XFS to

perform allocation and deallocation operations concurrently as long as concurrent operations affect different allocation groups. The number of concurrent operations that can be performed in the file system is therefore equal to the number of allocation groups. However, since the ability to perform concurrent operations is also limited by the number of processors able to perform the operations, Red Hat recommends that the number of allocation groups be greater than or equal to the number of processors in the system.

A single directory cannot be modified by multiple allocation groups simultaneously. Therefore, Red Hat recommends that applications that create and remove large numbers of files do not store all files in a single directory.

To configure allocation groups, use the **mkfs.xfs -d** option. See the **mkfs.xfs** man page for details.

Growth constraints

If you may need to increase the size of your file system after formatting time (either by adding more hardware or through thin-provisioning), you must carefully consider initial file layout, as allocation group size cannot be changed after formatting is complete.

Allocation groups must be sized according to the eventual capacity of the file system, not the initial capacity. The number of allocation groups in the fully-grown file system should not exceed several hundred, unless allocation groups are at their maximum size (1 TB). Therefore for most file systems, the recommended maximum growth to allow for a file system is ten times the initial size.

Additional care must be taken when growing a file system on a RAID array, as the device size must be aligned to an exact multiple of the allocation group size so that new allocation group headers are correctly aligned on the newly added storage. The new storage must also have the same geometry as the existing storage, since geometry cannot be changed after formatting time, and therefore cannot be optimized for storage of a different geometry on the same block device.

Inode size and inline attributes

If the inode has sufficient space available, XFS can write attribute names and values directly into the inode. These inline attributes can be retrieved and modified up to an order of magnitude faster than retrieving separate attribute blocks, as additional I/O is not required.

The default inode size is 256 bytes. Only around 100 bytes of this is available for attribute storage, depending on the number of data extent pointers stored in the inode. Increasing inode size when you format the file system can increase the amount of space available for storing attributes.

Both attribute names and attribute values are limited to a maximum size of 254 bytes. If either name or value exceeds 254 bytes in length, the attribute is pushed to a separate attribute block instead of being stored inline.

To configure inode parameters, use the **mkfs.xfs -i** option. See the **mkfs.xfs** man page for details.

RAID

If software RAID is in use, **mkfs.xfs** automatically configures the underlying hardware with an appropriate stripe unit and width. However, stripe unit and width may need to be manually configured if hardware RAID is in use, as not all hardware RAID devices export this information. To configure stripe unit and width, use the **mkfs.xfs -d** option. See the **mkfs.xfs** man page for details.

Log size

Pending changes are aggregated in memory until a synchronization event is triggered, at which point

they are written to the log. The size of the log determines the number of concurrent modifications that can be in-progress at one time. It also determines the maximum amount of change that can be aggregated in memory, and therefore how often logged data is written to disk. A smaller log forces data to be written back to disk more frequently than a larger log. However, a larger log uses more memory to record pending modifications, so a system with limited memory will not benefit from a larger log.

Logs perform better when they are aligned to the underlying stripe unit; that is, they start and end at stripe unit boundaries. To align logs to the stripe unit, use the **mkfs.xfs -d** option. See the **mkfs.xfs** man page for details.

To configure the log size, use the following **mkfs.xfs** option, replacing *logsize* with the size of the log:

```
# mkfs.xfs -l size=logsize
```

For further details, see the **mkfs.xfs** man page:

```
$ man mkfs.xfs
```

Log stripe unit

Log writes on storage devices that use RAID5 or RAID6 layouts may perform better when they start and end at stripe unit boundaries (are aligned to the underlying stripe unit). **mkfs.xfs** attempts to set an appropriate log stripe unit automatically, but this depends on the RAID device exporting this information.

Setting a large log stripe unit can harm performance if your workload triggers synchronization events very frequently, because smaller writes need to be padded to the size of the log stripe unit, which can increase latency. If your workload is bound by log write latency, Red Hat recommends setting the log stripe unit to 1 block so that unaligned log writes are triggered as possible.

The maximum supported log stripe unit is the size of the maximum log buffer size (256 KB). It is therefore possible that the underlying storage may have a larger stripe unit than can be configured on the log. In this case, **mkfs.xfs** issues a warning and sets a log stripe unit of 32 KB.

To configure the log stripe unit, use one of the following options, where *N* is the number of blocks to use as the stripe unit, and *size* is the size of the stripe unit in KB.

```
mkfs.xfs -l sunit=Nb
mkfs.xfs -l su=size
```

For further details, see the **mkfs.xfs** man page:

```
$ man mkfs.xfs
```

7.4.7.1.2. Mount Options

Inode allocation

- compress=zlib – Better compression ratio. It is the default and safe for older kernels.
- compress=lzo – Faster compression but does not compress as much as zlib.

- compress=no – Disables compression.
- compress-force=method – Enable compression even for files that do not compress well, such as videos and images of disks. The options are compress-force=zlib and compress-force=lzo.

Highly recommended for file systems greater than 1 TB in size. The ***inode64*** parameter configures XFS to allocate inodes and data across the entire file system. This ensures that inodes are not allocated largely at the beginning of the file system, and data is not largely allocated at the end of the file system, improving performance on large file systems.

Log buffer size and number

The larger the log buffer, the fewer I/O operations it takes to write all changes to the log. A larger log buffer can improve performance on systems with I/O-intensive workloads that do not have a non-volatile write cache.

The log buffer size is configured with the ***logbsize*** mount option, and defines the maximum amount of information that can be stored in the log buffer; if a log stripe unit is not set, buffer writes can be shorter than the maximum, and therefore there is no need to reduce the log buffer size for synchronization-heavy workloads. The default size of the log buffer is 32 KB. The maximum size is 256 KB and other supported sizes are 64 KB, 128 KB or power of 2 multiples of the log stripe unit between 32 KB and 256 KB.

The number of log buffers is defined by the ***logbufs*** mount option. The default value is 8 log buffers (the maximum), but as few as two log buffers can be configured. It is usually not necessary to reduce the number of log buffers, except on memory-bound systems that cannot afford to allocate memory to additional log buffers. Reducing the number of log buffers tends to reduce log performance, especially on workloads sensitive to log I/O latency.

Delay change logging

XFS has the option to aggregate changes in memory before writing them to the log. The ***delaylog*** parameter allows frequently modified metadata to be written to the log periodically instead of every time it changes. This option increases the potential number of operations lost in a crash and increases the amount of memory used to track metadata. However, it can also increase metadata modification speed and scalability by an order of magnitude, and does not reduce data or metadata integrity when **fsync**, **fdatasync**, or **sync** are used to ensure data and metadata is written to disk.

7.4.7.2. Tuning ext4

This section covers some of the tuning parameters available to ext4 file systems at format and at mount time.

7.4.7.2.1. Formatting Options

Inode table initialization

Initializing all inodes in the file system can take a very long time on very large file systems. By default, the initialization process is deferred (lazy inode table initialization is enabled). However, if your system does not have an ext4 driver, lazy inode table initialization is disabled by default. It can be enabled by setting **lazy_itable_init** to 1). In this case, kernel processes continue to initialize the file system after it is mounted.

This section describes only some of the options available at format time. For further formatting parameters, see the **mkfs.ext4** man page:

```
$ man mkfs.ext4
```

7.4.7.2.2. Mount Options

Inode table initialization rate

When lazy inode table initialization is enabled, you can control the rate at which initialization occurs by specifying a value for the ***init_itable*** parameter. The amount of time spent performing background initialization is approximately equal to 1 divided by the value of this parameter. The default value is **10**.

Automatic file synchronization

Some applications do not correctly perform an **fsync** after renaming an existing file, or after truncating and rewriting. By default, ext4 automatically synchronizes files after each of these operations. However, this can be time consuming.

If this level of synchronization is not required, you can disable this behavior by specifying the **noauto_da_alloc** option at mount time. If **noauto_da_alloc** is set, applications must explicitly use fsync to ensure data persistence.

Journal I/O priority

By default, journal I/O has a priority of **3**, which is slightly higher than the priority of normal I/O. You can control the priority of journal I/O with the ***journal_ioprio*** parameter at mount time. Valid values for ***journal_ioprio*** range from **0** to **7**, with **0** being the highest priority I/O.

This section describes only some of the options available at mount time. For further mount options, see the **mount** man page:

```
$ man mount
```

7.4.7.3. Tuning Btrfs

Starting with Red Hat Enterprise Linux 7.0, Btrfs is provided as a Technology Preview. Tuning should always be done to optimize the system based on its current workload. For information on creation and mounting options, see the chapter on Btrfs in the Red Hat Enterprise Linux 7 Storage Administration Guide.

Data Compression

The default compression algorithm is zlib, but a specific workload can give a reason to change the compression algorithm. For example, if you have a single thread with heavy file I/O, using the lzo algorithm can be more preferable. Options at mount time are:

- **compress=zlib** – the default option with a high compression ratio, safe for older kernels.
- **compress=lzo** – compression faster, but lower, than zlib.
- **compress=no** – disables compression.
- **compress-force=*method*** – enables compression even for files that do not compress well, such as videos and disk images. The available methods are **zlib** and **lzo**.

Only files created or changed after the mount option is added will be compressed. To compress existing files, run the following command after you replace *method* with either **zlib** or **lzo**:

```
$ btrfs filesystem defragment -cmethod
```

To re-compress the file using **lzo**, run:

```
$ btrfs filesystem defragment -r -v -clzo /
```

7.4.7.4. Tuning GFS2

This section covers some of the tuning parameters available to GFS2 file systems at format and at mount time.

Directory spacing

All directories created in the top-level directory of the GFS2 mount point are automatically spaced to reduce fragmentation and increase write speed in those directories. To space another directory like a top-level directory, mark that directory with the ***T*** attribute, as shown, replacing *dirname* with the path to the directory you wish to space:

```
# chattr +T dirname
```

chattr is provided as part of the e2fsprogs package.

Reduce contention

GFS2 uses a global locking mechanism that can require communication between the nodes of a cluster. Contention for files and directories between multiple nodes lowers performance. You can minimize the risk of cross-cache invalidation by minimizing the areas of the file system that are shared between multiple nodes.

CHAPTER 8. NETWORKING

The networking subsystem is comprised of a number of different parts with sensitive connections. Red Hat Enterprise Linux 7 networking is therefore designed to provide optimal performance for most workloads, and to optimize its performance automatically. As such, it is not usually necessary to manually tune network performance. This chapter discusses further optimizations that can be made to functional networking systems.

Network performance problems are sometimes the result of hardware malfunction or faulty infrastructure. Resolving these issues is beyond the scope of this document.

8.1. CONSIDERATIONS

To make good tuning decisions, you need a thorough understanding of packet reception in Red Hat Enterprise Linux. This section explains how network packets are received and processed, and where potential bottlenecks can occur.

A packet sent to a Red Hat Enterprise Linux system is received by the network interface card (NIC) and placed in either an internal hardware buffer or a ring buffer. The NIC then sends a hardware interrupt request, prompting the creation of a software interrupt operation to handle the interrupt request.

As part of the software interrupt operation, the packet is transferred from the buffer to the network stack. Depending on the packet and your network configuration, the packet is then forwarded, discarded, or passed to a socket receive queue for an application and then removed from the network stack. This process continues until either there are no packets left in the NIC hardware buffer, or a certain number of packets (specified in **`/proc/sys/net/core/dev_weight`**) are transferred.

The Red Hat Enterprise Linux Network Performance Tuning Guide available on the Red Hat Customer Portal contains information on packet reception in the Linux kernel, and covers the following areas of NIC tuning: SoftIRQ misses (netdev budget), **`tuned`** tuning daemon, **`numad`** NUMA daemon, CPU power states, interrupt balancing, pause frames, interrupt coalescence, adapter queue (**`netdev`** backlog), adapter RX and TX buffers, adapter TX queue, module parameters, adapter offloading, Jumbo Frames, TCP and UDP protocol tuning, and NUMA locality.

8.1.1. Before You Tune

Network performance problems are most often the result of hardware malfunction or faulty infrastructure. Red Hat highly recommends verifying that your hardware and infrastructure are working as expected before beginning to tune the network stack.

8.1.2. Bottlenecks in Packet Reception

While the network stack is largely self-optimizing, there are a number of points during network packet processing that can become bottlenecks and reduce performance.

The NIC hardware buffer or ring buffer

The hardware buffer might be a bottleneck if a large number of packets are being dropped. For information about monitoring your system for dropped packets, see Section 8.2.4, “ethtool”.

The hardware or software interrupt queues

Interrupts can increase latency and processor contention. For information on how interrupts are handled by the processor, see Section 5.1.3, “Interrupt Request (IRQ) Handling”. For information on how to monitor interrupt handling in your system, see Section 5.2.3, “/proc/interrupts”. For configuration options that affect interrupt handling, see Section 5.3.7, “Setting Interrupt Affinity”.

The socket receive queue for the application

A bottleneck in an application's receive queue is indicated by a large number of packets that are not copied to the requesting application, or by an increase in UDP input errors (**`InErrors`**) in **`/proc/net/snmp`**. For information about monitoring your system for these errors, see Section 8.2.1, “ss” and Section 8.2.5, “/proc/net/snmp”.

8.2. MONITORING AND DIAGNOSING PERFORMANCE PROBLEMS

Red Hat Enterprise Linux 7 provides a number of tools that are useful for monitoring system performance and diagnosing performance problems related to the networking subsystem. This section outlines the available tools and gives examples of how to use them to monitor and diagnose network related performance issues.

8.2.1. ss

ss is a command-line utility that prints statistical information about sockets, allowing administrators to assess device performance over time. By default, **ss** lists open non-listening TCP sockets that have established connections, but a number of useful options are provided to help administrators filter out statistics about specific sockets.

Red Hat recommends **ss** over **netstat** in Red Hat Enterprise Linux 7.

ss is provided by the iproute package. For more information, see the man page:

```
$ man ss
```

8.2.2. ip

The **ip** utility lets administrators manage and monitor routes, devices, routing policies, and tunnels. The **`ip monitor`** command can continuously monitor the state of devices, addresses, and routes.

ip is provided by the iproute package. For details about using ip, see the man page:

```
$ man ip
```

8.2.3. dropwatch

Dropwatch is an interactive tool that monitors and records packets that are dropped by the kernel.

For further information, see the **dropwatch** man page:

```
$ man dropwatch
```

8.2.4. ethtool

The **ethtool** utility allows administrators to view and edit network interface card settings. It is useful for observing the statistics of certain devices, such as the number of packets dropped by that device.

You can view the status of a specified device's counters with **`ethtool -S`** and the name of the device you want to monitor.

```
$ ethtool -S devname
```

For further information, see the man page:

```
$ man ethtool
```

8.2.5. /proc/net/snmp

The **/proc/net/snmp** file displays data that is used by snmp agents for IP, ICMP, TCP and UDP monitoring and management. Examining this file on a regular basis can help administrators identify unusual values and thereby identify potential performance problems. For example, an increase in UDP input errors (**InErrors**) in **/proc/net/snmp** can indicate a bottleneck in a socket receive queue.

8.2.6. Network Monitoring with SystemTap

The *Red Hat Enterprise Linux 7 SystemTap Beginner's Guide* includes several sample scripts that are useful for profiling and monitoring network performance.

The following **SystemTap** example scripts relate to networking and may be useful in diagnosing network performance problems. By default they are installed to the **/usr/share/doc/systemtap-client/examples/network** directory.

nettop.stp

Every 5 seconds, prints a list of processes (process identifier and command) with the number of packets sent and received and the amount of data sent and received by the process during that interval.

socket-trace.stp

Instruments each of the functions in the Linux kernel's **net/socket.c** file, and prints trace data.

dropwatch.stp

Every 5 seconds, prints the number of socket buffers freed at locations in the kernel. Use the **--all-modules** option to see symbolic names.

The **latencytap.stp** script records the effect that different types of latency have on one or more processes. It prints a list of latency types every 30 seconds, sorted in descending order by the total time the process or processes spent waiting. This can be useful for identifying the cause of both storage and network latency. Red Hat recommends using the **--all-modules** option with this script to better enable the mapping of latency events. By default, this script is installed to the **/usr/share/doc/systemtap-client-*version*/examples/profiling** directory.

For further information, see the *Red Hat Enterprise Linux 7 SystemTap Beginner's Guide*, available from http://access.redhat.com/site/documentation/Red_Hat_Enterprise_Linux/.

8.3. CONFIGURATION TOOLS

Red Hat Enterprise Linux provides a number of tools to assist administrators in configuring the system. This section outlines the available tools and provides examples of how they can be used to solve network related performance problems in Red Hat Enterprise Linux 7.

However, it is important to keep in mind that network performance problems are sometimes the result of hardware malfunction or faulty infrastructure. Red Hat highly recommends verifying that your hardware and infrastructure are working as expected before using these tools to tune the network stack.

Further, some network performance problems are better resolved by altering the application than by reconfiguring your network subsystem. It is generally a good idea to configure your application to perform frequent posix calls, even if this means queuing data in the application space, as this allows data to be stored flexibly and swapped in or out of memory as required.

8.3.1. Tuned Profiles for Network Performance

The **Tuned** service provides a number of different profiles to improve performance in a number of specific use cases. The following profiles can be useful for improving networking performance.

- latency-performance
- network-latency
- network-throughput

For more information about these profiles, see Section A.5, “tuned-adm”.

8.3.2. Configuring the Hardware Buffer

If a large number of packets are being dropped by the hardware buffer, there are a number of potential solutions.

Slow the input traffic

Filter incoming traffic, reduce the number of joined multicast groups, or reduce the amount of broadcast traffic to decrease the rate at which the queue fills. For details of how to filter incoming traffic, see the *Red Hat Enterprise Linux 7 Security Guide*. For details about multicast groups, see the Red Hat Enterprise Linux 7 Clustering documentation. For details about broadcast traffic, see the *Red Hat Enterprise Linux 7 System Administrator's Reference Guide*, or documentation related to the device you want to configure. All Red Hat Enterprise Linux 7 documentation is available from http://access.redhat.com/site/documentation/Red_Hat_Enterprise_Linux/.

Resize the hardware buffer queue

Reduce the number of packets being dropped by increasing the size of the queue so that the it does not overflow as easily. You can modify the rx/tx parameters of the network device with the ethtool command:

```
# ethtool --set-ring devname value
```

Change the drain rate of the queue

Device weight refers to the number of packets a device can receive at one time (in a single scheduled processor access). You can increase the rate at which a queue is drained by increasing its device weight, which is controlled by the ***dev_weight*** parameter. This parameter can be temporarily altered by changing the contents of the **/proc/sys/net/core/dev_weight** file, or permanently altered with **sysctl**, which is provided by the procps-ng package.

Altering the drain rate of a queue is usually the simplest way to mitigate poor network performance. However, increasing the number of packets that a device can receive at one time uses additional processor time, during which no other processes can be scheduled, so this can cause other performance

problems.

8.3.3. Configuring Interrupt Queues

If analysis reveals high latency, your system may benefit from poll-based rather than interrupt-based packet receipt.

8.3.3.1. Configuring Busy Polling

Busy polling helps reduce latency in the network receive path by allowing socket layer code to poll the receive queue of a network device, and disabling network interrupts. This removes delays caused by the interrupt and the resultant context switch. However, it also increases CPU utilization. Busy polling also prevents the CPU from sleeping, which can incur additional power consumption.

Busy polling is disabled by default. To enable busy polling on specific sockets, do the following.

- Set ***sysctl.net.core.busy_poll*** to a value other than **0**. This parameter controls the number of microseconds to wait for packets on the device queue for socket poll and selects. Red Hat recommends a value of **50**.
- Add the **SO_BUSY_POLL** socket option to the socket.

To enable busy polling globally, you must also set ***sysctl.net.core.busy_read*** to a value other than **0**. This parameter controls the number of microseconds to wait for packets on the device queue for socket reads. It also sets the default value of the **SO_BUSY_POLL** option. Red Hat recommends a value of **50** for a small number of sockets, and a value of **100** for large numbers of sockets. For extremely large numbers of sockets (more than several hundred), use **epoll** instead.

Busy polling behavior is supported by the following drivers. These drivers are also supported on Red Hat Enterprise Linux 7.

- bnx2x
- be2net
- ixgbe
- mlx4
- myri10ge

As of Red Hat Enterprise Linux 7.1, you can also run the following command to check whether a specific device supports busy polling.

```
# ethtool -k device | grep "busy-poll"
```

If this returns **busy-poll: on [fixed]**, busy polling is available on the device.

8.3.4. Configuring Socket Receive Queues

If analysis suggests that packets are being dropped because the drain rate of a socket queue is too slow, there are several ways to alleviate the performance issues that result.

Decrease the speed of incoming traffic

Decrease the rate at which the queue fills by filtering or dropping packets before they reach the queue, or by lowering the weight of the device.

Increase the depth of the application's socket queue

If a socket queue that receives a limited amount of traffic in bursts, increasing the depth of the socket queue to match the size of the bursts of traffic may prevent packets from being dropped.

8.3.4.1. Decrease the Speed of Incoming Traffic

Filter incoming traffic or lower the network interface card's device weight to slow incoming traffic. For details of how to filter incoming traffic, see the *Red Hat Enterprise Linux 7 Security Guide*, available from http://access.redhat.com/site/documentation/Red_Hat_Enterprise_Linux/.

Device weight refers to the number of packets a device can receive at one time (in a single scheduled processor access). Device weight is controlled by the ***dev_weight*** parameter. This parameter can be temporarily altered by changing the contents of the **/proc/sys/net/core/dev_weight** file, or permanently altered with **sysctl**, which is provided by the procps-ng package.

8.3.4.2. Increasing Queue Depth

Increasing the depth of an application socket queue is typically the easiest way to improve the drain rate of a socket queue, but it is unlikely to be a long-term solution.

To increase the depth of a queue, increase the size of the socket receive buffer by making either of the following changes:

Increase the value of /proc/sys/net/core/rmem_default

This parameter controls the default size of the receive buffer used by sockets. This value must be smaller than or equal to the value of **/proc/sys/net/core/rmem_max**.

Use setsockopt to configure a larger SO_RCVBUF value

This parameter controls the maximum size in bytes of a socket's receive buffer. Use the **getsockopt** system call to determine the current value of the buffer. For further information, see the socket(7) manual page.

8.3.5. Configuring Receive-Side Scaling (RSS)

Receive-Side Scaling (RSS), also known as multi-queue receive, distributes network receive processing across several hardware-based receive queues, allowing inbound network traffic to be processed by multiple CPUs. RSS can be used to relieve bottlenecks in receive interrupt processing caused by overloading a single CPU, and to reduce network latency.

To determine whether your network interface card supports RSS, check whether multiple interrupt request queues are associated with the interface in **/proc/interrupts**. For example, if you are interested in the **p1p1** interface:

```
# egrep 'CPU|p1p1' /proc/interrupts
   CPU0    CPU1    CPU2    CPU3    CPU4    CPU5
89:   40187       0       0       0       0       0   IR-PCI-MSI-edge
p1p1-0
90:       0     790       0       0       0       0   IR-PCI-MSI-edge
```

```
p1p1-1
91:       0       0     959       0       0       0   IR-PCI-MSI-edge
p1p1-2
92:       0       0       0    3310       0       0   IR-PCI-MSI-edge
p1p1-3
93:       0       0       0       0     622       0   IR-PCI-MSI-edge
p1p1-4
94:       0       0       0       0       0    2475   IR-PCI-MSI-edge
p1p1-5
```

The preceding output shows that the NIC driver created 6 receive queues for the **p1p1** interface (**p1p1-0** through **p1p1-5**). It also shows how many interrupts were processed by each queue, and which CPU serviced the interrupt. In this case, there are 6 queues because by default, this particular NIC driver creates one queue per CPU, and this system has 6 CPUs. This is a fairly common pattern amongst NIC drivers.

Alternatively, you can check the output of **ls -1 /sys/devices/*/*/*device_pci_address*/msi_irqs** after the network driver is loaded. For example, if you are interested in a device with a PCI address of **0000:01:00.0**, you can list the interrupt request queues of that device with the following command:

```
# ls -1 /sys/devices/*/*/0000:01:00.0/msi_irqs
101
102
103
104
105
106
107
108
109
```

RSS is enabled by default. The number of queues (or the CPUs that should process network activity) for RSS are configured in the appropriate network device driver. For the **bnx2x** driver, it is configured in ***num_queues***. For the **sfc** driver, it is configured in the ***rss_cpus*** parameter. Regardless, it is typically configured in **/sys/class/net/*device*/queues/*rx-queue*/**, where *device* is the name of the network device (such as **eth1**) and *rx-queue* is the name of the appropriate receive queue.

When configuring RSS, Red Hat recommends limiting the number of queues to one per physical CPU core. Hyper-threads are often represented as separate cores in analysis tools, but configuring queues for all cores including logical cores such as hyper-threads has not proven beneficial to network performance.

When enabled, RSS distributes network processing equally between available CPUs based on the amount of processing each CPU has queued. However, you can use the **ethtool *--show-rxfh-indir*** and ***--set-rxfh-indir*** parameters to modify how network activity is distributed, and weight certain types of network activity as more important than others.

The **irqbalance** daemon can be used in conjunction with RSS to reduce the likelihood of cross-node memory transfers and cache line bouncing. This lowers the latency of processing network packets.

8.3.6. Configuring Receive Packet Steering (RPS)

Receive Packet Steering (RPS) is similar to RSS in that it is used to direct packets to specific CPUs for processing. However, RPS is implemented at the software level, and helps to prevent the hardware queue of a single network interface card from becoming a bottleneck in network traffic.

RPS has several advantages over hardware-based RSS:

- RPS can be used with any network interface card.
- It is easy to add software filters to RPS to deal with new protocols.
- RPS does not increase the hardware interrupt rate of the network device. However, it does introduce inter-processor interrupts.

RPS is configured per network device and receive queue, in the **/sys/class/net/*device*/queues/*rx-queue*/rps_cpus** file, where *device* is the name of the network device (such as **eth0**) and *rx-queue* is the name of the appropriate receive queue (such as **rx-0**).

The default value of the **rps_cpus** file is **0**. This disables RPS, so the CPU that handles the network interrupt also processes the packet.

To enable RPS, configure the appropriate **rps_cpus** file with the CPUs that should process packets from the specified network device and receive queue.

The **rps_cpus** files use comma-delimited CPU bitmaps. Therefore, to allow a CPU to handle interrupts for the receive queue on an interface, set the value of their positions in the bitmap to 1. For example, to handle interrupts with CPUs 0, 1, 2, and 3, set the value of **rps_cpus** to **f**, which is the hexadecimal value for 15. In binary representation, 15 is **00001111** (1+2+4+8).

For network devices with single transmit queues, best performance can be achieved by configuring RPS to use CPUs in the same memory domain. On non-NUMA systems, this means that all available CPUs can be used. If the network interrupt rate is extremely high, excluding the CPU that handles network interrupts may also improve performance.

For network devices with multiple queues, there is typically no benefit to configuring both RPS and RSS, as RSS is configured to map a CPU to each receive queue by default. However, RPS may still be beneficial if there are fewer hardware queues than CPUs, and RPS is configured to use CPUs in the same memory domain.

8.3.7. Configuring Receive Flow Steering (RFS)

Receive Flow Steering (RFS) extends RPS behavior to increase the CPU cache hit rate and thereby reduce network latency. Where RPS forwards packets based solely on queue length, RFS uses the RPS backend to calculate the most appropriate CPU, then forwards packets based on the location of the application consuming the packet. This increases CPU cache efficiency.

RFS is disabled by default. To enable RFS, you must edit two files:

/proc/sys/net/core/rps_sock_flow_entries

Set the value of this file to the maximum expected number of concurrently active connections. We recommend a value of **32768** for moderate server loads. All values entered are rounded up to the nearest power of 2 in practice.

/sys/class/net/*device*/queues/*rx-queue*/rps_flow_cnt

Replace *device* with the name of the network device you wish to configure (for example, **eth0**), and *rx-queue* with the receive queue you wish to configure (for example, **rx-0**).

Set the value of this file to the value of **rps_sock_flow_entries** divided by **N**, where **N** is the number of receive queues on a device. For example, if **rps_flow_entries** is set to **32768** and there are 16 configured receive queues, **rps_flow_cnt** should be set to **2048**. For single-queue devices, the value of **rps_flow_cnt** is the same as the value of **rps_sock_flow_entries**.

Data received from a single sender is not sent to more than one CPU. If the amount of data received from a single sender is greater than a single CPU can handle, configure a larger frame size to reduce the number of interrupts and therefore the amount of processing work for the CPU. Alternatively, consider NIC offload options or faster CPUs.

Consider using **numactl** or **taskset** in conjunction with RFS to pin applications to specific cores, sockets, or NUMA nodes. This can help prevent packets from being processed out of order.

8.3.8. Configuring Accelerated RFS

Accelerated RFS boosts the speed of RFS by adding hardware assistance. Like RFS, packets are forwarded based on the location of the application consuming the packet. Unlike traditional RFS, however, packets are sent directly to a CPU that is local to the thread consuming the data: either the CPU that is executing the application, or a CPU local to that CPU in the cache hierarchy.

Accelerated RFS is only available if the following conditions are met:

- Accelerated RFS must be supported by the network interface card. Accelerated RFS is supported by cards that export the **ndo_rx_flow_steer() netdevice** function.
- **ntuple** filtering must be enabled.

Once these conditions are met, CPU to queue mapping is deduced automatically based on traditional RFS configuration. That is, CPU to queue mapping is deduced based on the IRQ affinities configured by the driver for each receive queue. Refer to Section 8.3.7, “Configuring Receive Flow Steering (RFS)” for details on configuring traditional RFS.

Red Hat recommends using accelerated RFS wherever using RFS is appropriate and the network interface card supports hardware acceleration.

APPENDIX A. TOOL REFERENCE

This appendix provides a quick reference for the various tools in Red Hat Enterprise Linux 7 that can be used to tweak performance. See the relevant man page for your tool for complete, up-to-date, detailed reference material.

A.1. IRQBALANCE

irqbalance is a command line tool that distributes hardware interrupts across processors to improve system performance. It runs as a daemon by default, but can be run once only with the **`--oneshot`** option.

The following parameters are useful for improving performance.

--powerthresh

Sets the number of CPUs that can idle before a CPU is placed into powersave mode. If more CPUs than the threshold are more than 1 standard deviation below the average softirq workload and no CPUs are more than one standard deviation above the average, and have more than one irq assigned to them, a CPU is placed into powersave mode. In powersave mode, a CPU is not part of irq balancing so that it is not woken unnecessarily.

--hintpolicy

Determines how irq kernel affinity hinting is handled. Valid values are **`exact`** (irq affinity hint is always applied), **`subset`** (irq is balanced, but the assigned object is a subset of the affinity hint), or **`ignore`** (irq affinity hint is ignored completely).

--policyscript

Defines the location of a script to execute for each interrupt request, with the device path and irq number passed as arguments, and a zero exit code expected by **irqbalance**. The script defined can specify zero or more key value pairs to guide **irqbalance** in managing the passed irq.

The following are recognized as valid key value pairs.

ban

Valid values are **`true`** (exclude the passed irq from balancing) or **`false`** (perform balancing on this irq).

balance_level

Allows user override of the balance level of the passed irq. By default the balance level is based on the PCI device class of the device that owns the irq. Valid values are **none**, **`package`**, **`cache`**, or **`core`**.

numa_node

Allows user override of the NUMA node that is considered local to the passed irq. If information about the local node is not specified in ACPI, devices are considered equidistant from all nodes. Valid values are integers (starting from 0) that identify a specific NUMA node, and **`-1`**, which specifies that an irq should be considered equidistant from all nodes.

--banirq

The interrupt with the specified interrupt request number is added to the list of banned interrupts.

You can also use the ***IRQBALANCE_BANNED_CPUS*** environment variable to specify a mask of CPUs that are ignored by **irqbalance**.

For further details, see the man page:

```
$ man irqbalance
```

A.2. ETHTOOL

The **ethtool** utility allows administrators to view and edit network interface card settings. It is useful for observing the statistics of certain devices, such as the number of packets dropped by that device.

ethtool, its options, and its usage, are comprehensively documented on the man page.

```
$ man ethtool
```

A.3. SS

ss is a command-line utility that prints statistical information about sockets, allowing administrators to assess device performance over time. By default, **ss** lists open non-listening TCP sockets that have established connections, but a number of useful options are provided to help administrators filter out statistics about specific sockets.

One commonly used command is **`ss -tmpie`**, which displays all TCP sockets (**`t`**, internal TCP information (**`i`**), socket memory usage (**`m`**), processes using the socket (**`p`**), and detailed socket information (**`i`**).

Red Hat recommends **ss** over **netstat** in Red Hat Enterprise Linux 7.

ss is provided by the iproute package. For more information, see the man page:

```
$ man ss
```

A.4. TUNED

Tuned is a tuning daemon that can adapt the operating system to perform better under certain workloads by setting a tuning profile. It can also be configured to react to changes in CPU and network use and adjusts settings to improve performance in active devices and reduce power consumption in inactive devices.

To configure dynamic tuning behavior, edit the ***`dynamic_tuning`*** parameter in the **`/etc/tuned/tuned-main.conf`** file. **Tuned** then periodically analyzes system statistics and uses them to update your system tuning settings. You can configure the time interval in seconds between these updates with the ***`update_interval`*** parameter.

For further details about tuned, see the man page:

```
$ man tuned
```

A.5. TUNED-ADM

tuned-adm is a command line tool that enables you to switch between **Tuned** profiles to improve performance in a number of specific use cases. It also provides the **`tuned-adm recommend`** sub-command that assesses your system and outputs a recommended tuning profile.

As of Red Hat Enterprise Linux 7, **Tuned** includes the ability to run any shell command as part of enabling or disabling a tuning profile. This allows you to extend **Tuned** profiles with functionality that has not been integrated into **Tuned** yet.

Red Hat Enterprise Linux 7 also provides the ***`include`*** parameter in profile definition files, allowing you to base your own **Tuned** profiles on existing profiles.

The following tuning profiles are provided with **Tuned** and are supported in Red Hat Enterprise Linux 7.

throughput-performance

A server profile focused on improving throughput. This is the default profile, and is recommended for most systems.

This profile favors performance over power savings by setting **`intel_pstate`** and **`min_perf_pct=100`**. It enables transparent huge pages, uses **cpupower** to set the **`performance`** cpufreq governor, and sets the input/output scheduler to **`deadline`**. It also sets ***`kernel.sched_min_granularity_ns`*** to **`10`** µs, ***`kernel.sched_wakeup_granularity_ns`*** to **`15`** µs, and ***`vm.dirty_ratio`*** to **`40`**%.

latency-performance

A server profile focused on lowering latency. This profile is recommended for latency-sensitive workloads that benefit from c-state tuning and the increased TLB efficiency of transparent huge pages.

This profile favors performance over power savings by setting **`intel_pstate`** and **`max_perf_pct=100`**. It enables transparent huge pages, uses **cpupower** to set the **`performance`** cpufreq governor, and requests a ***`cpu_dma_latency`*** value of **`1`**.

network-latency

A server profile focused on lowering network latency.

This profile favors performance over power savings by setting **`intel_pstate`** and **`min_perf_pct=100`**. It disables transparent huge pages, and automatic NUMA balancing. It also uses **cpupower** to set the **`performance`** cpufreq governor, and requests a ***`cpu_dma_latency`*** value of **`1`**. It also sets ***`busy_read`*** and ***`busy_poll`*** times to **`50`** µs, and ***`tcp_fastopen`*** to **`3`**.

network-throughput

A server profile focused on improving network throughput.

This profile favors performance over power savings by setting **`intel_pstate`** and **`max_perf_pct=100`** and increasing kernel network buffer sizes. It enables transparent huge pages, and uses **cpupower** to set the **`performance`** cpufreq governor. It also sets ***`kernel.sched_min_granularity_ns`*** to **`10`** µs, ***`kernel.sched_wakeup_granularity_ns`*** to 15 µs, and ***`vm.dirty_ratio`*** to **`40`**%.

virtual-guest

A profile focused on optimizing performance in Red Hat Enterprise Linux 7 virtual machines as well as VMware guests.

This profile favors performance over power savings by setting **intel_pstate** and **max_perf_pct=100**. It also decreases the swappiness of virtual memory. It enables transparent huge pages, and uses **cpupower** to set the **performance** cpufreq governor. It also sets ***kernel.sched_min_granularity_ns*** to **10** µs, ***kernel.sched_wakeup_granularity_ns*** to 15 µs, and ***vm.dirty_ratio*** to **40**%.

virtual-host

A profile focused on optimizing performance in Red Hat Enterprise Linux 7 virtualization hosts.

This profile favors performance over power savings by setting **intel_pstate** and **max_perf_pct=100**. It also decreases the swappiness of virtual memory. This profile enables transparent huge pages and writes dirty pages back to disk more frequently. It uses **cpupower** to set the **performance** cpufreq governor. It also sets ***kernel.sched_min_granularity_ns*** to **10** µs, ***kernel.sched_wakeup_granularity_ns*** to 15 µs, ***kernel.sched_migration_cost*** to **5** µs, and ***vm.dirty_ratio*** to **40**%.

For detailed information about the power saving profiles provided with tuned-adm, see the *Red Hat Enterprise Linux 7 Power Management Guide*, available from http://access.redhat.com/site/documentation/Red_Hat_Enterprise_Linux/.

For detailed information about using **tuned-adm**, see the man page:

```
$ man tuned-adm
```

A.6. PERF

The **perf** tool provides a number of useful commands, some of which are listed in this section. For detailed information about **perf**, see the *Red Hat Enterprise Linux 7 Developer Guide*, available from http://access.redhat.com/site/documentation/Red_Hat_Enterprise_Linux/, or refer to the man pages.

perf stat

This command provides overall statistics for common performance events, including instructions executed and clock cycles consumed. You can use the option flags to gather statistics on events other than the default measurement events. As of Red Hat Enterprise Linux 6.4, it is possible to use **perf stat** to filter monitoring based on one or more specified control groups (cgroups).

For further information, read the man page:

```
$ man perf-stat
```

perf record

This command records performance data into a file which can be later analyzed using **perf report**. For further details, read the man page:

```
$ man perf-record
```

perf report

This command reads the performance data from a file and analyzes the recorded data. For further details, read the man page:

```
$ man perf-report
```

perf list

This command lists the events available on a particular machine. These events vary based on the performance monitoring hardware and the software configuration of the system. For further information, read the man page:

```
$ man perf-list
```

perf top

This command performs a similar function to the **top** tool. It generates and displays a performance counter profile in realtime. For further information, read the man page:

```
$ man perf-top
```

perf trace

This command performs a similar function to the **strace** tool. It monitors the system calls used by a specified thread or process and all signals received by that application. Additional trace targets are available; refer to the man page for a full list:

```
$ man perf-trace
```

A.7. PERFORMANCE CO-PILOT (PCP)

Performance Co-Pilot (PCP) provides a large number of command-line tools, graphical tools, and libraries. For more information on these tools, see their respective manual pages.

Table A.1. System Services Distributed with Performance Co-Pilot in Red Hat Enterprise Linux 7

Name	Description
pmcd	The Performance Co-Pilot Collection Daemon (PMCD).
pmie	The Performance Metrics Inference Engine.
pmlogger	The performance metrics logger.
pmmgr	Manages a collection of PCP daemons for a set of discovered local and remote hosts running the Performance Co-Pilot Collection Daemon (PMCD) according to zero or more configuration directories.
pmproxy	The Performance Co-Pilot Collection Daemon (PMCD) proxy server.
pmwebd	Binds a subset of the Performance Co-Pilot client API to RESTful web applications using the HTTP protocol.

Table A.2. Tools Distributed with Performance Co-Pilot in Red Hat Enterprise Linux 7

Name	Description
pcp	Displays the current status of a Performance Co-Pilot installation.
pmatop	Shows the system-level occupation of the most critical hardware resources from the performance point of view: CPU, memory, disk, and network.
pmchart	Plots performance metrics values available through the facilities of the Performance Co-Pilot.
pmclient	Displays high-level system performance metrics by using the Performance Metrics Application Programming Interface (PMAPI).
pmcollectl	Collects and displays system-level data, either from a live system or from a Performance Co-Pilot archive file.
pmconfig	Displays the values of configuration parameters.
pmdbg	Displays available Performance Co-Pilot debug control flags and their values.
pmdiff	Compares the average values for every metric in either one or two archives, in a given time window, for changes that are likely to be of interest when searching for performance regressions.
pmdumplog	Displays control, metadata, index, and state information from a Performance Co-Pilot archive file.
pmdumptext	Outputs the values of performance metrics collected live or from a Performance Co-Pilot archive.
pmerr	Displays available Performance Co-Pilot error codes and their corresponding error messages.
pmfind	Finds PCP services on the network.
pmie	An inference engine that periodically evaluates a set of arithmetic, logical, and rule expressions. The metrics are collected either from a live system, or from a Performance Co-Pilot archive file.
pmieconf	Displays or sets configurable **pmie** variables.
pminfo	Displays information about performance metrics. The metrics are collected either from a live system, or from a Performance Co-Pilot archive file.
pmiostat	Reports I/O statistics for SCSI devices (by default) or **device-mapper** devices (with the **-x dm** option).
pmlc	Interactively configures active **pmlogger** instances.

Name	Description
pmlogcheck	Identifies invalid data in a Performance Co-Pilot archive file.
pmlogconf	Creates and modifies a **pmlogger** configuration file.
pmloglabel	Verifies, modifies, or repairs the label of a Performance Co-Pilot archive file.
pmlogsummary	Calculates statistical information about performance metrics stored in a Performance Co-Pilot archive file.
pmprobe	Determines the availability of performance metrics.
pmrep	Reports on selected, easily customizable, performance metrics values.
pmsocks	Allows access to a Performance Co-Pilot hosts through a firewall.
pmstat	Periodically displays a brief summary of system performance.
pmstore	Modifies the values of performance metrics.
pmtrace	Provides a command line interface to the trace Performance Metrics Domain Agent (PMDA).
pmval	Displays the current value of a performance metric.

A.8. VMSTAT

Vmstat outputs reports on your system's processes, memory, paging, block input/output, interrupts, and CPU activity. It provides an instantaneous report of the average of these events since the machine was last booted, or since the previous report.

-a

Displays active and inactive memory.

-f

Displays the number of forks since boot. This includes the **fork**, **vfork**, and **clone** system calls, and is equivalent to the total number of tasks created. Each process is represented by one or more tasks, depending on thread usage. This display does not repeat.

-m

Displays slab information.

-n

Specifies that the header will appear once, not periodically.

-s

Displays a table of various event counters and memory statistics. This display does not repeat.

delay

The delay between reports in seconds. If no delay is specified, only one report is printed, with the average values since the machine was last booted.

count

The number of times to report on the system. If no count is specified and delay is defined, **vmstat** reports indefinitely.

-d

Displays disk statistics.

-p

Takes a partition name as a value, and reports detailed statistics for that partition.

-S

Defines the units output by the report. Valid values are **k** (1000 bytes), **K** (1024 bytes), **m** (1,000,000 bytes), or **M** (1,048,576 bytes).

-D

Report summary statistics about disk activity.

For detailed information about the output provided by each output mode, see the man page:

```
$ man vmstat
```

A.9. X86_ENERGY_PERF_POLICY

The **x86_energy_perf_policy** tool allows administrators to define the relative importance of performance and energy efficiency. It is provided by the kernel-tools package.

To view the current policy, run the following command:

```
# x86_energy_perf_policy -r
```

To set a new policy, run the following command:

```
# x86_energy_perf_policy profile_name
```

Replace profile_name with one of the following profiles.

performance

The processor does not sacrifice performance for the sake of saving energy. This is the default value.

normal

The processor tolerates minor performance compromises for potentially significant energy savings. This is a reasonable saving for most servers and desktops.

powersave

The processor accepts potentially significant performance decreases in order to maximize energy efficiency.

For further details of how to use **x86_energy_perf_policy**, see the man page:

```
$ man x86_energy_perf_policy
```

A.10. TURBOSTAT

The **turbostat** tool provides detailed information about the amount of time that the system spends in different states. **Turbostat** is provided by the kernel-tools package.

By default, **turbostat** prints a summary of counter results for the entire system, followed by counter results every 5 seconds, under the following headings:

pkg

The processor package number.

core

The processor core number.

CPU

The Linux CPU (logical processor) number.

%c0

The percentage of the interval for which the CPU retired instructions.

GHz

When this number is higher than the value in TSC, the CPU is in turbo mode

TSC

The average clock speed over the course of the entire interval.

%c1, %c3, and %c6

The percentage of the interval for which the processor was in the c1, c3, or c6 state, respectively.

%pc3 or %pc6

The percentage of the interval for which the processor was in the pc3 or pc6 state, respectively.

Specify a different period between counter results with the **-i** option, for example, run **turbostat -i 10** to print results every 10 seconds instead.

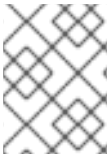

NOTE

Upcoming Intel processors may add additional c-states. As of Red Hat Enterprise Linux 7.0, **turbostat** provides support for the c7, c8, c9, and c10 states.

A.11. NUMASTAT

The **numastat** tool is provided by the numactl package, and displays memory statistics (such as allocation hits and misses) for processes and the operating system on a per-NUMA-node basis. The default tracking categories for the **numastat** command are outlined as follows:

numa_hit

The number of pages that were successfully allocated to this node.

numa_miss

The number of pages that were allocated on this node because of low memory on the intended node. Each **numa_miss** event has a corresponding **numa_foreign** event on another node.

numa_foreign

The number of pages initially intended for this node that were allocated to another node instead. Each **numa_foreign** event has a corresponding **numa_miss** event on another node.

interleave_hit

The number of interleave policy pages successfully allocated to this node.

local_node

The number of pages successfully allocated on this node, by a process on this node.

other_node

The number of pages allocated on this node, by a process on another node.

Supplying any of the following options changes the displayed units to megabytes of memory (rounded to two decimal places), and changes other specific **numastat** behaviors as described below.

-c

Horizontally condenses the displayed table of information. This is useful on systems with a large number of NUMA nodes, but column width and inter-column spacing are somewhat unpredictable. When this option is used, the amount of memory is rounded to the nearest megabyte.

-m

Displays system-wide memory usage information on a per-node basis, similar to the information found in **/proc/meminfo**.

-n

Displays the same information as the original numastat command (**numa_hit**, **numa_miss**, **numa_foreign**, **interleave_hit**, **local_node**, and **other_node**), with an updated format, using megabytes as the unit of measurement.

-p pattern

Displays per-node memory information for the specified pattern. If the value for pattern is comprised of digits, **numastat** assumes that it is a numerical process identifier. Otherwise, **numastat** searches process command lines for the specified pattern.

Command line arguments entered after the value of the **-p** option are assumed to be additional patterns for which to filter. Additional patterns expand, rather than narrow, the filter.

-s

Sorts the displayed data in descending order so that the biggest memory consumers (according to the total column) are listed first.

Optionally, you can specify a node, and the table will be sorted according to the node column. When using this option, the node value must follow the **-s** option immediately, as shown here:

```
numastat -s2
```

Do not include white space between the option and its value.

-v

Displays more verbose information. Namely, process information for multiple processes will display detailed information for each process.

-V

Displays numastat version information.

-z

Omits table rows and columns with only zero values from the displayed information. Note that some near-zero values that are rounded to zero for display purposes will not be omitted from the displayed output.

A.12. NUMACTL

Numactl lets administrators run a process with a specified scheduling or memory placement policy. **Numactl** can also set a persistent policy for shared memory segments or files, and set the processor affinity and memory affinity of a process.

Numactl provides a number of useful options. This appendix outlines some of these options and gives suggestions for their use, but is not exhaustive.

--hardware

Displays an inventory of available nodes on the system, including relative distances between nodes.

--membind

Ensures that memory is allocated only from specific nodes. If there is insufficient memory available in the specified location, allocation fails.

--cpunodebind

Ensures that a specified command and its child processes execute only on the specified node.

--phycpubind

Ensures that a specified command and its child processes execute only on the specified processor.

--localalloc

Specifies that memory should always be allocated from the local node.

--preferred

Specifies a preferred node from which to allocate memory. If memory cannot be allocated from this specified node, another node will be used as a fallback.

For further details about these and other parameters, see the man page:

```
$ man numactl
```

A.13. NUMAD

numad is an automatic NUMA affinity management daemon. It monitors NUMA topology and resource usage within a system in order to dynamically improve NUMA resource allocation and management.

Note that when **numad** is enabled, its behavior overrides the default behavior of automatic NUMA balancing.

A.13.1. Using numad from the Command Line

To use **numad** as an executable, just run:

```
# numad
```

While **numad** runs, its activities are logged in **/var/log/numad.log**. It will run until stopped with the following command:

```
# numad -i 0
```

Stopping **numad** does not remove the changes it has made to improve NUMA affinity. If system use changes significantly, running **numad** again will adjust affinity to improve performance under the new conditions.

To restrict **numad** management to a specific process, start it with the following options.

```
# numad -S 0 -p pid
```

-p *pid*

This option adds the specified *pid* to an explicit inclusion list. The process specified will not be managed until it meets the **numad** process significance threshold.

-S 0

This sets the type of process scanning to **0**, which limits **numad** management to explicitly included processes.

For further information about available **numad** options, refer to the **numad** man page:

```
$ man numad
```

A.13.2. Using numad as a Service

While **numad** runs as a service, it attempts to tune the system dynamically based on the current system workload. Its activities are logged in **/var/log/numad.log**.

To start the service, run:

```
# systemctl start numad.service
```

To make the service persist across reboots, run:

```
# chkconfig numad on
```

For further information about available **numad** options, refer to the **numad** man page:

```
$ man numad
```

A.13.3. Pre-Placement Advice

numad provides a pre-placement advice service that can be queried by various job management systems to provide assistance with the initial binding of CPU and memory resources for their processes. This pre-placement advice is available regardless of whether **numad** is running as an executable or a service.

A.13.4. Using numad with KSM

If KSM is in use on a NUMA system, change the value of the **/sys/kernel/mm/ksm/merge_nodes** parameter to **0** to avoid merging pages across NUMA nodes. Otherwise, KSM increases remote memory accesses as it merges pages across nodes. Furthermore, kernel memory accounting statistics can eventually contradict each other after large amounts of cross-node merging. As such, **numad** can become confused about the correct amounts and locations of available memory, after the KSM daemon merges many memory pages. KSM is beneficial only if you are overcommitting the memory on your system. If your system has sufficient free memory, you may achieve higher performance by turning off and disabling the KSM daemon.

A.14. OPROFILE

OProfile is a low overhead, system-wide performance monitoring tool provided by the oprofile package. It uses the performance monitoring hardware on the processor to retrieve information about the kernel and executables on the system, such as when memory is referenced, the number of second-level cache requests, and the number of hardware interrupts received. OProfile is also able to profile applications that run in a Java Virtual Machine (JVM).

OProfile provides the following tools. Note that the legacy **opcontrol** tool and the new **operf** tool are mutually exclusive.

ophelp

Displays available events for the system's processor along with a brief description of each.

opimport

Converts sample database files from a foreign binary format to the native format for the system. Only use this option when analyzing a sample database from a different architecture.

opannotate

Creates annotated source for an executable if the application was compiled with debugging symbols.

opcontrol

Configures which data is collected in a profiling run.

operf

Intended to replace **opcontrol**. The **operf** tool uses the Linux Performance Events subsystem, allowing you to target your profiling more precisely, as a single process or system-wide, and allowing OProfile to co-exist better with other tools using the performance monitoring hardware on your system. Unlike **opcontrol**, no initial setup is required, and it can be used without the root privileges unless the **--system-wide** option is in use.

opreport

Retrieves profile data.

oprofiled

Runs as a daemon to periodically write sample data to disk.

Legacy mode (**opcontrol**, **oprofiled**, and post-processing tools) remains available, but is no longer the recommended profiling method.

For further information about any of these commands, see the OProfile man page:

```
$ man oprofile
```

A.15. TASKSET

The **taskset** tool is provided by the util-linux package. It allows administrators to retrieve and set the processor affinity of a running process, or launch a process with a specified processor affinity.

IMPORTANT

taskset does not guarantee local memory allocation. If you require the additional performance benefits of local memory allocation, Red Hat recommends using **numactl** instead of taskset.

To set the CPU affinity of a running process, run the following command:

```
# taskset -c processors pid
```

Replace *processors* with a comma delimited list of processors or ranges of processors (for example, **1,3,5-7**. Replace *pid* with the process identifier of the process that you want to reconfigure.

To launch a process with a specified affinity, run the following command:

```
# taskset -c processors -- application
```

Replace *processors* with a comma delimited list of processors or ranges of processors. Replace *application* with the command, options and arguments of the application you want to run.

For more information about **taskset**, see the man page:

```
$ man taskset
```

A.16. SYSTEMTAP

SystemTap is extensively documented in its own guides. The Red Hat Enterprise Linux 7 versions of the *SystemTap Beginner's Guide* and the *SystemTap TapSet Reference* are both available from http://access.redhat.com/site/documentation/Red_Hat_Enterprise_Linux/.

APPENDIX B. REVISION HISTORY

Revision 10.13-55 **Mon Nov 6 2017** **Marek Suchánek**
Asynchronous update.

Revision 10.13-50 **Thu Jul 27 2017** **Milan Navrátil**
Document version for 7.4 GA publication.

Revision 10.13-44 **Tue Dec 13 2016** **Milan Navrátil**
Asynchronous update.

Revision 10.08-38 **Wed Nov 11 2015** **Jana Heves**
Version for 7.2 GA release.

Revision 0.3-23 **Tue Feb 17 2015** **Laura Bailey**
Building for RHEL 7.1 GA.

Revision 0.3-3 **Mon Apr 07 2014** **Laura Bailey**
Rebuilding for RHEL 7.0 GA.